



# Faculty of Engineering & IT

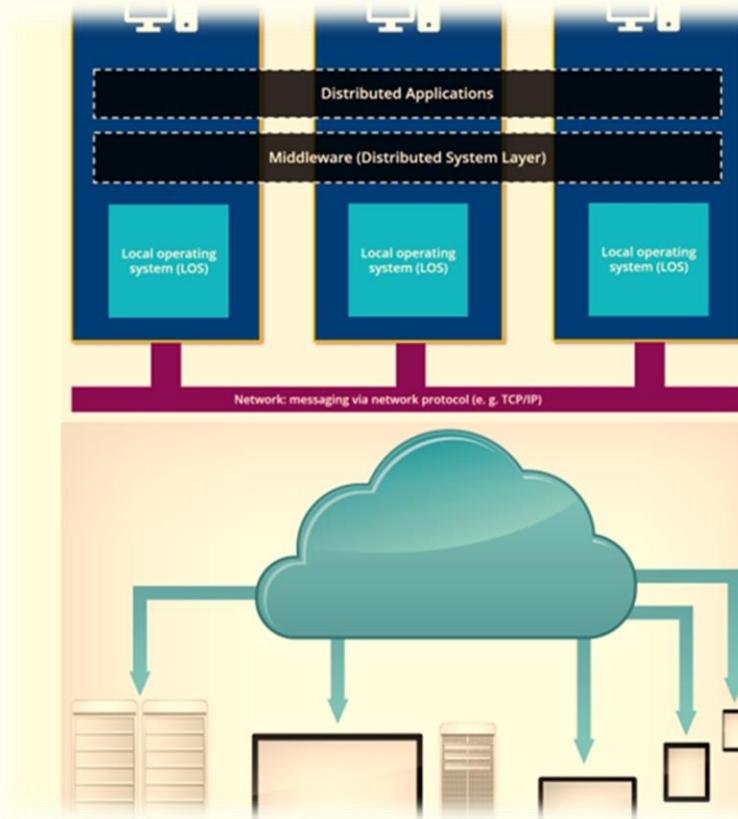
## Software Engineering Department

# ITSE5336 – DISTRIBUTED & CLOUD COMPUTING



*Dr Mosab M. Ayiad*

Spring 2022/2023



Demonstrating distributed systems  
architectures, middleware, & cloud computing

# MODULE OVERVIEW

# Module Description:

---

- In the mid of 1980s, microprocessors and high-speed networking were the two main advances that changed software systems. Approaches for parallel and distributed computing were introduced and continually evolved.
- Nowadays, powerful mini-computers (e.g., smartphones) are inter-networked and provide innovative services enabled by pervasive and cloud technology.
- This module takes students on an educational journey that begins with the concepts and models of distributed computing and systems and proceeds toward the paradigms of pervasive and cloud computing.
- The module covers several topics, ranging from the architecture of distributed systems, design issues, and common major problems to the recent cloud technologies, layouts, and services.
- The module is arranged into an incremental approach and the learning material is introduced in a gradual manner. The teaching involves practice occasions with chosen learning items for students to enhance their comprehension and skills.
- This module does not cover Security issues of distributed and cloud computing.

# Aims and Objectives:

---

- The module aims to:
  1. Introduce students to distributed and cloud systems.
  2. Empower the understanding of the difference of systems architecture and application development.
  3. Introduce students to cloud computing environments, challenges and services.
- Module objectives:
  1. Understand the design challenges and implementation problems in distributed computing.
  2. Equip students with the knowledge to solve and handle distributed computing issues.
  3. Enable students to practice the concepts of parallel and distributed systems programming.
  4. Understand the design patterns of cloud applications.
  5. Demonstrate recent implementation issues in cloud computing.

# Module Outcomes:

---

- By the end of this module students:

1. Understand various system models such as: layered-based, event-based and object-based architectures.
2. Understand and able to implement parallel processes and distributed applications.
3. Being able to design and implement systems using remote procedure call and message passing techniques.
4. Being able to identify and solve distribution design issues such as: naming, agreement, scalability, availability and reliability.
5. Understand the layouts of cloud computing and main could environments.
6. Being able to design and implement cloud systems using centralised and decentralised approaches.
7. Understand the design and implementation issues raised by ubiquitous computing and the engagement of IoT.

# Module Structure:

---

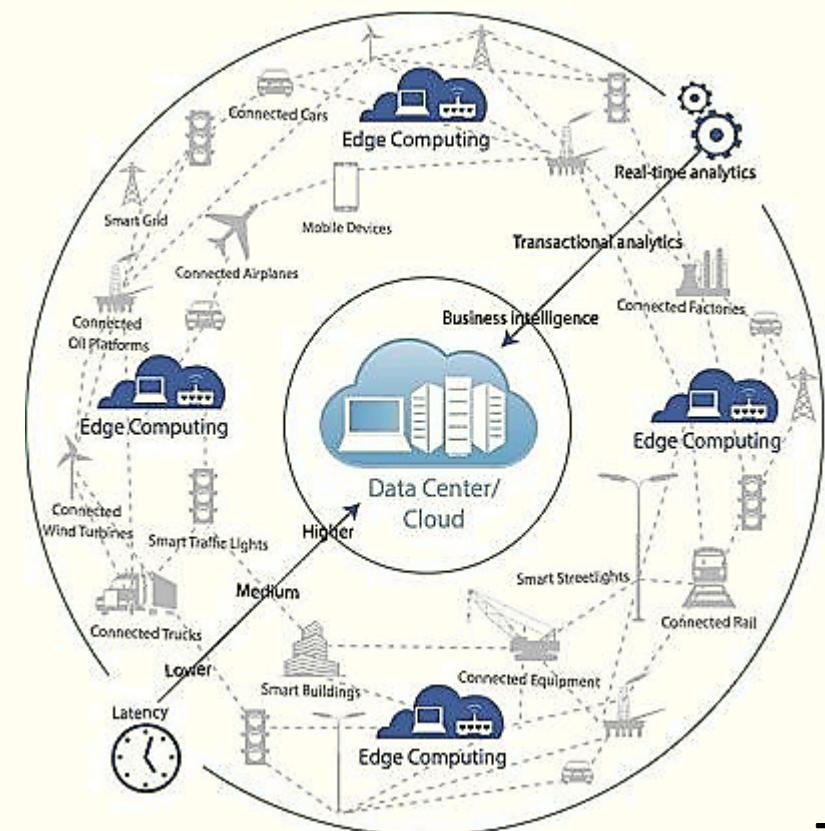
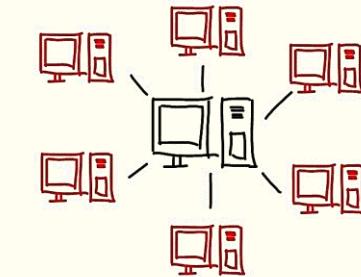
- The lectures involve:
  - Frontal lectures, paper presentations, in-class exercises, and team works.
  - The first portion of the module introduces basic concepts and approaches in distributed and cloud computing.
  - The next portion focuses on the concepts in distributed computing; middleware and major issues and solution approaches.
  - The final portion considers cloud technology. It involves the main cloud providers, models, and novel services.

# Module Syllabus:

## Weeks:

- 01:** Module overview.
- 02, 03:** Introduction.
- 04, 05:** Architectures of Distributed Computing (DC).
- 06:** Processes.
- 07:** Communications.
- 08:** Reading & mid-term examinations.
- 09:** Reading & mid-term examinations.
- 10:** Naming & Coordination.
- 11:** Replication & Consistency & Fault-tolerance.
- 12:** Orientation to Cloud Computing (CC).
- 13:** CC Environments.
- 14:** Virtualization in CC.
- 15:** Fog, Edge, & Ubiquitous Computing.
- 16:** Final examinations.

## Distributed Computing



## Practice sessions:

- There will be **no in-lab** practice sessions in this module, however, there are **coursework assignments**, **in-class practices**, and **homework**.

## Grading:

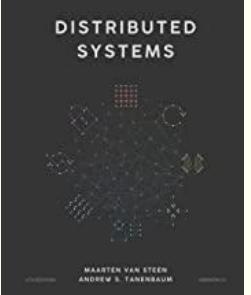
Pre-midterm assignments:	<b>20%</b> , to be submitted before Midterm exam (Week <b>3, 7</b> ).
Midterm exam:	<b>30%</b> , <b>~Week 8</b> .
Third assignment:	<b>10%</b> , to be submitted before Final exam (Week <b>13</b> ).
Final exam:	<b>40%</b> .

## Vital notes:

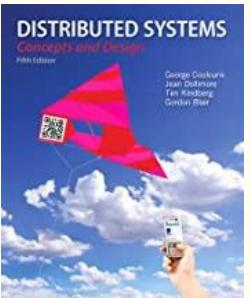
- ❑ All assignments are per-student tasks.
- ❑ The submission of all assignments will be via Moodle.
- ❑ All exams are conducted in invigilated paper examinations.

# Text Books and References:

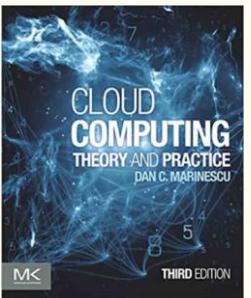
- I. M. van Steen, A. S. Tanenbaum, “*Distributed Systems*”, 4<sup>th</sup> ed, 2023, Pearson Education.



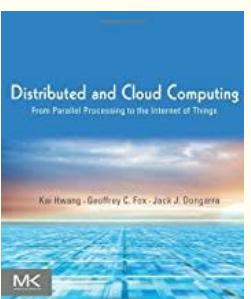
- II. G. Coulouris, J. Dollimore, T. Kindberg, G. Blair, “*Distributed Systems Concept and Design*”, 5<sup>th</sup> ed, 2012, Pearson.



- III. D. C. Marinescu, “*Cloud Computing Theory and Practice*”, 3<sup>rd</sup> ed, 2022, Morgan Kaufmann.



- IV. K. Hwang, G. Fox, J. Dongarra, “*Distributed and Cloud Computing from Parallel Processing to the Internet of Things*”, 1<sup>st</sup> ed, 2012, Morgan Kaufmann.



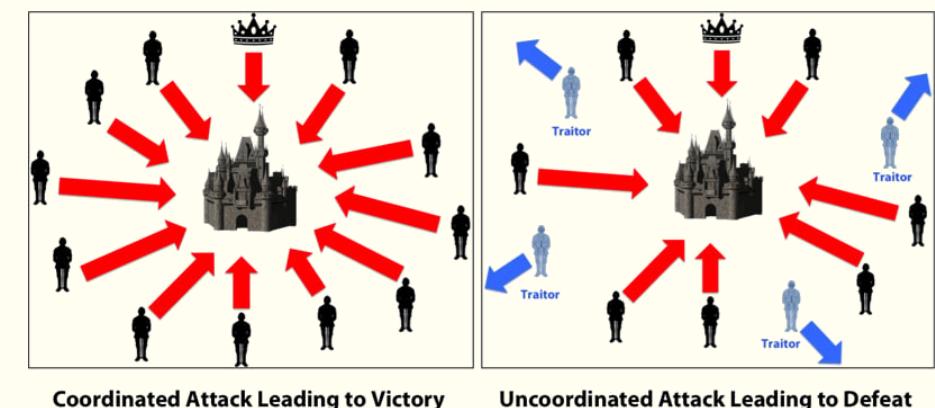
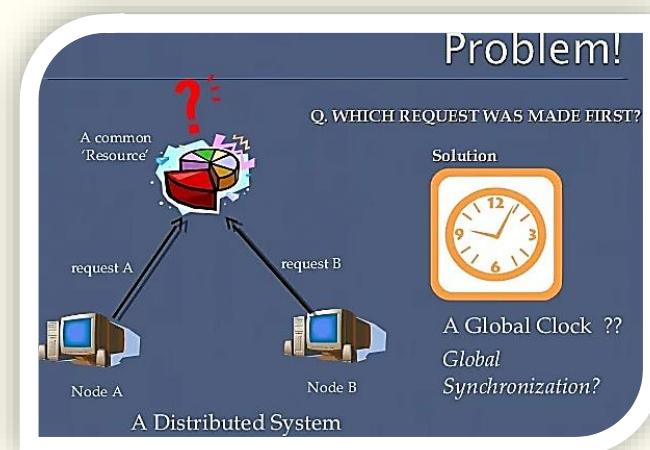
- Additional reading materials will be provided during the module.

# INTRODUCTION

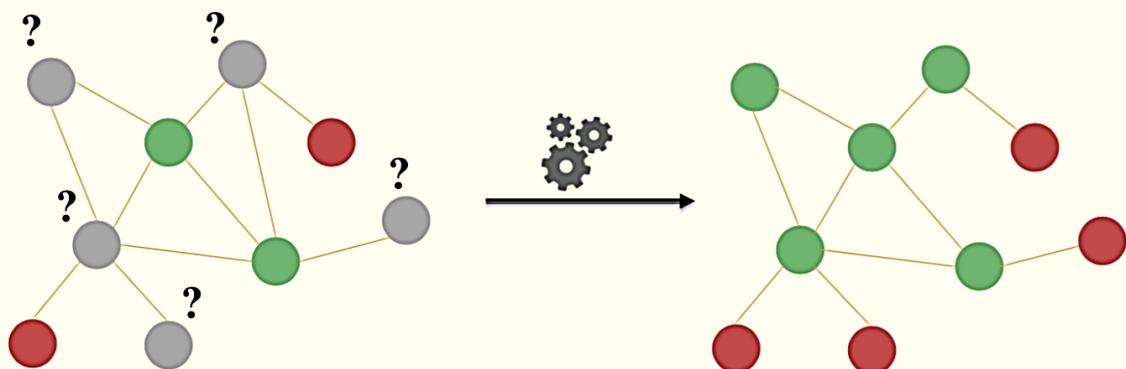
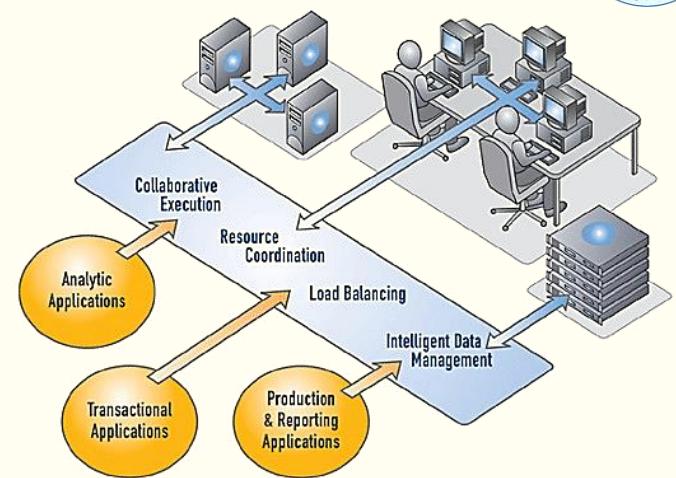
# Definitions:

- **Distributed System (DS)** “is a collection of autonomous computing elements that appears to its users as a **single coherent system**” Van Steen & Tanenbaum.

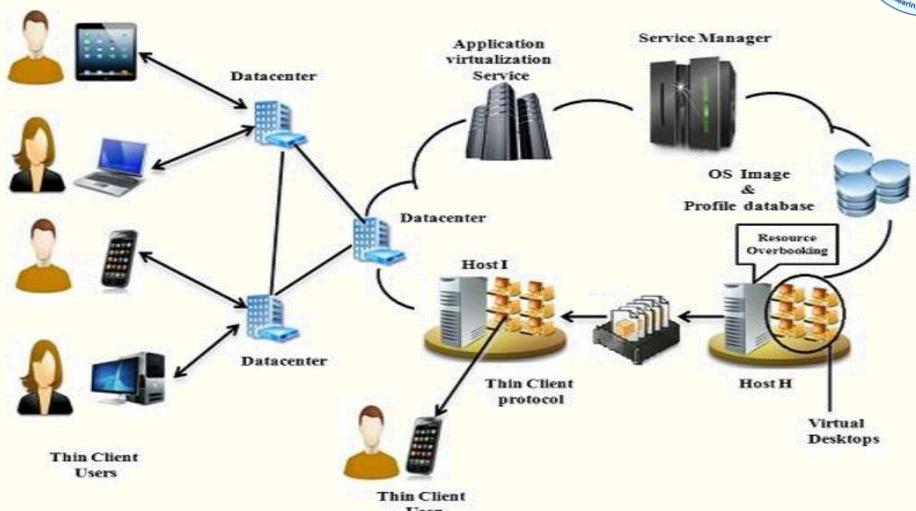
- Autonomous computing elements (a.k.a **nodes**), e.g., hardware **devices** or software **processes**.
  - **Preserve Independent behaviour**, i.e., each node is autonomous and will thus have its own notion of **time**: there is no **global clock**, which leads to fundamental **synchronization** and **coordination** problems.
- **Single coherent system**: users or applications perceive a single system  $\Rightarrow$  nodes need to **collaborate**.
  - All nodes operate the same. It doesn't matter where, when, and how interaction with the system takes place (**Transparency**), e.g., consider a banking system.
  - How to manage **membership**?
  - How to know that you are indeed communicating with an **authorized member**, e.g., consider the **Byzantine Failure**?



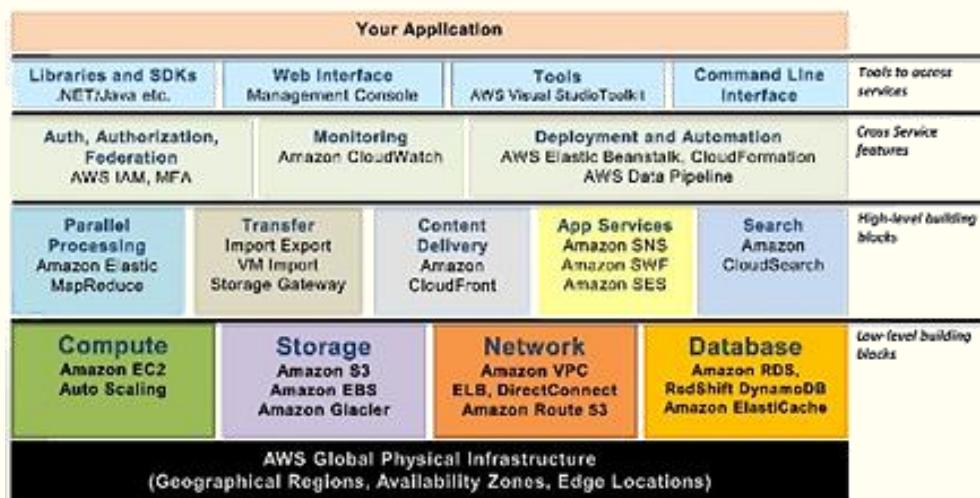
- Distributed Computing (DC) refers to a software that solves a problem over distributed autonomous computers which communicate over a network.
- DC aims to achieve computational tasks more faster than using a single compute.
- It is so characterised by Remote Procedure Calls (RPC) and Remote Method Invocation (RMI) computations.
  - Nodes collaborate by message-passing.
  - A common language is a requirement, e.g., CORBA, SOAP.
  - Nodes must trust each other.
- Classification of DC systems:
  - Distributed Information Systems.
  - Distributed Pervasive Systems.
  - Middleware.
- Limitations: reliability and scalability issues, nodes and networking may fail.



- Cloud Computing (CC) provides on demand **resources/services**, e.g., storage, database, analytics, software etc. over the **Internet!!!**
- CC implements various **techniques** to delivers hosted **services** over the Internet to cloud **users/customers**. It is characterised by providing a shared pool of **configurable computing resources, on-demand services, and pay-per-use services**, etc.
- Cloud systems are classified into different types, e.g.,
  - Public Cloud.
  - Private Cloud.
  - Community Cloud.
  - Hybrid Cloud.
- CC compromise some general requirements, e.g., being cost effective, fixable, reliable, scalable, globally accessible, secure, etc.
- Limitations: less control, security and privacy issues.

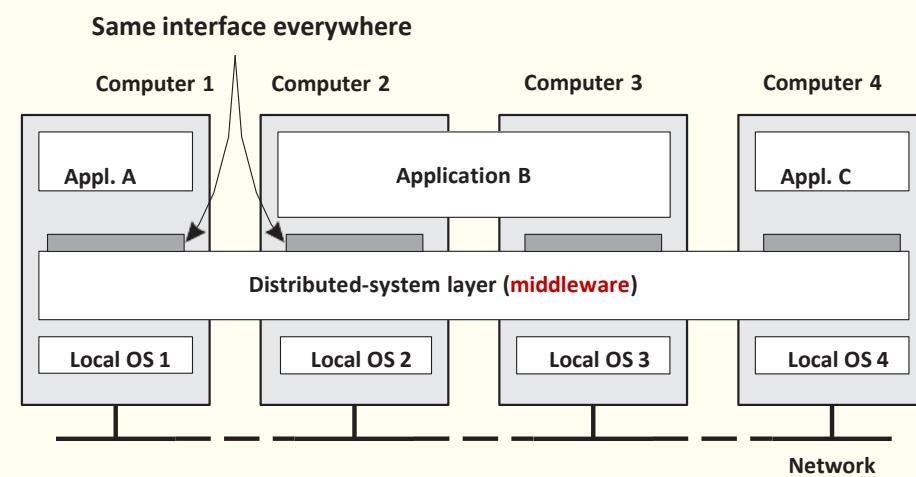
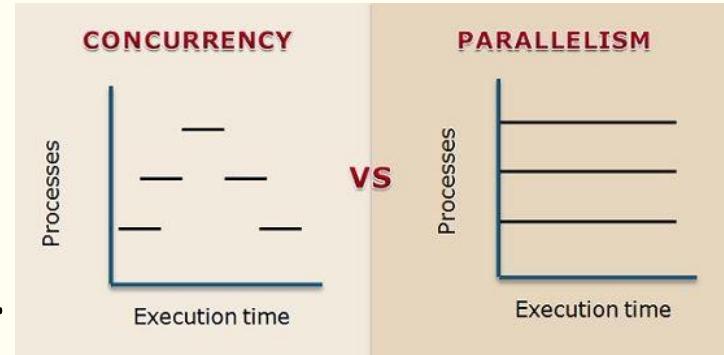


## AWS Cloud Layers



# Engineering of DS:

- Intentional requirements of DS (Why):
  - Be **cost-effective**, via **sharing data** and **computing resources**.
  - Be **high performance**, via **Parallelism & Concurrency**.
  - Be **scalable**, through **cloning of serving nodes (Replication)**.
  - Be **reliable**, by using **Redundancy & Fault-tolerance** techniques.
- Engineering requirements of DS (What & How):
  - Resource Sharing (Support **Global Shares**).
  - Transparency (Hide **Complexity**).
  - Openness (Use **Interoperable Components**).
  - Scalability (Enable **Extensibility**).
  - Availability (Ensure **99%** up-time).
  - Maintainability (Make it **easy** to **maintain & evolve**).
- DS requirements characterize the **complications** and **diversity** of the involved **techniques, mechanisms, and approaches**.



## • Resource Sharing:

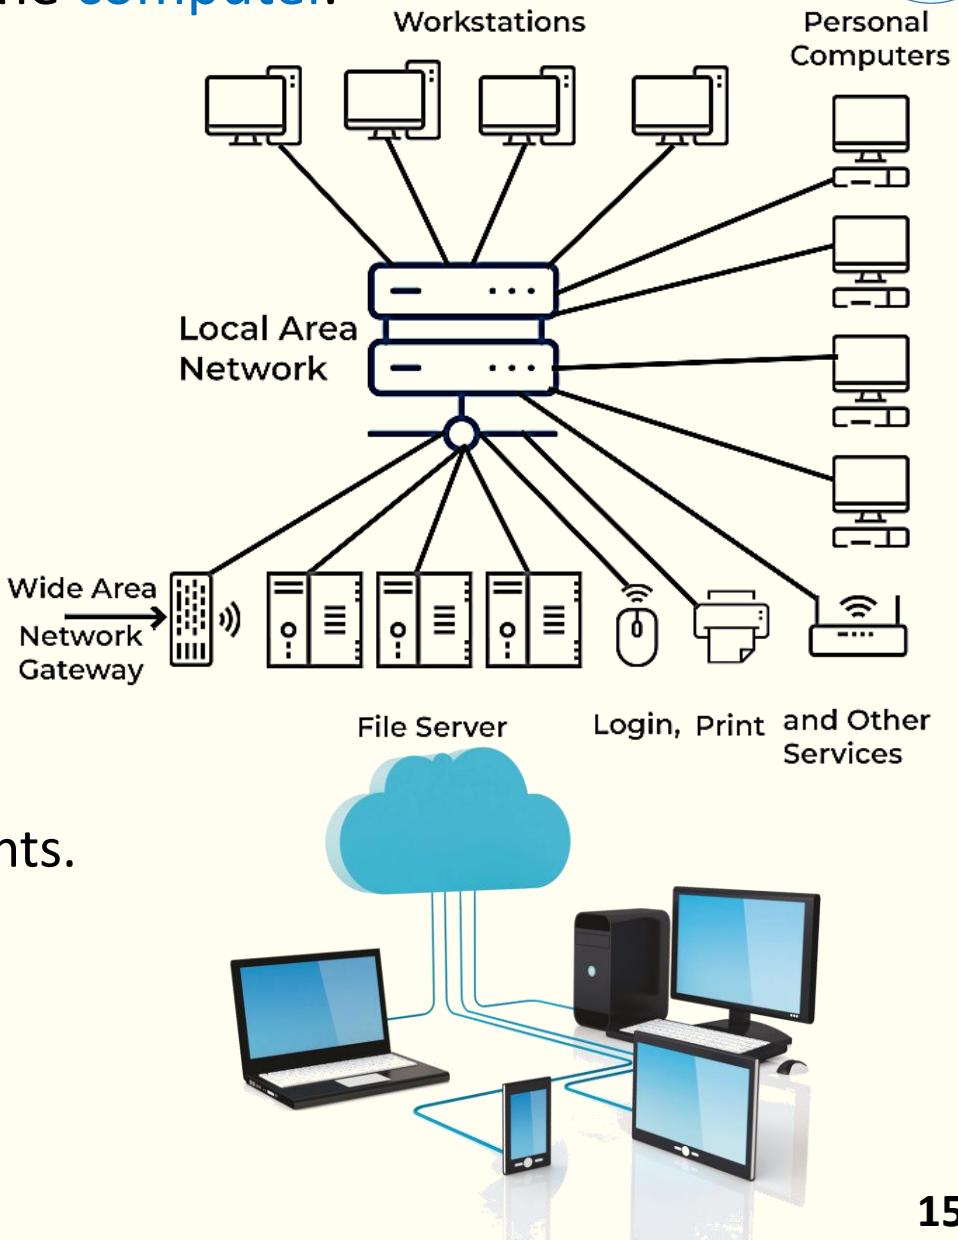
- With **stable** resource sharing, the **network** becomes the **computer**.

### • Examples:

- Internet connection, Printer, and File sharing.
- Database & Database servers.
- Peer-to-Peer (**P2P**) multimedia sharing.
- Mail and web hosting servers.
- Cloud computing and storage sharing.

### • Limitations:

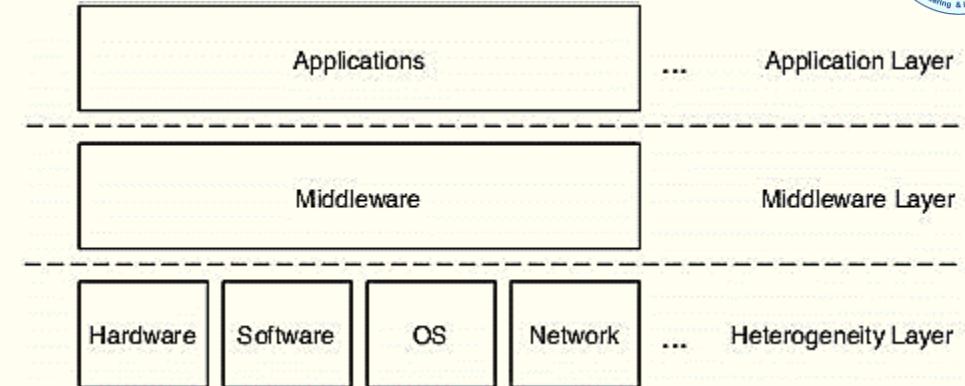
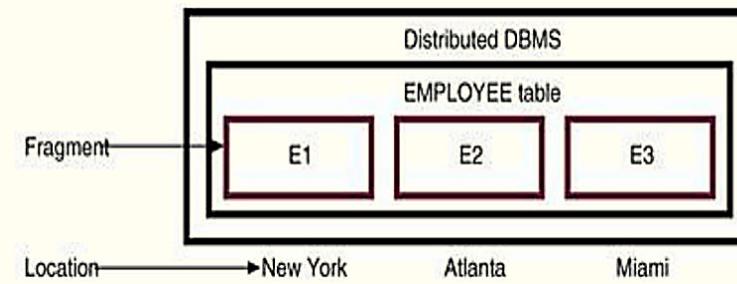
- Subject to local and remote **failures**.
- May experience high **latencies** or **delays**.
- Lack of **standardization & agreement** among participants.
- Availability** and **formats** issues.
- Security** and **authentication** issues.



## • Distribution Transparency:

- Examples:

- Distributed DBMS.
- Middleware.



- Transparency Types:

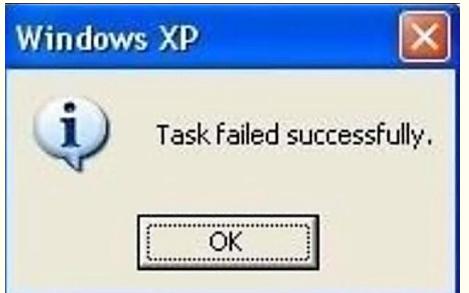
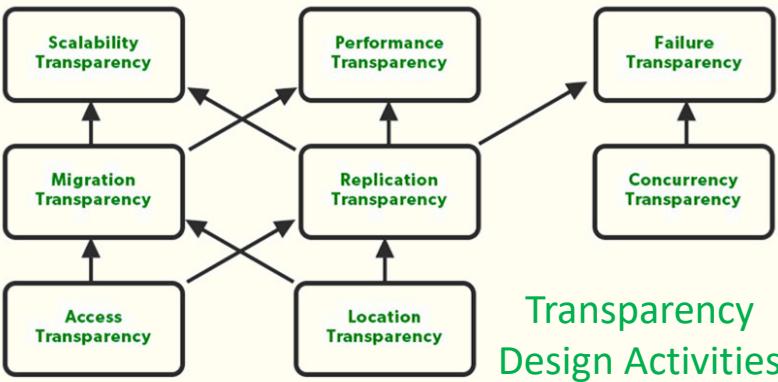
- **Access**: hide differences in data representation and how an object is accessed.
- **Location & Migration**: hide where an object is located; or hide that an object may be moved to another location while in use.
- **Replication**: use an object or its replica as same object.
- **Concurrency**: hide that an object may be shared by several independent users.
- **Failure**: hide the failure and recovery of an object.

- Follow the link for "[Transparency: Illusions of a Single System](#)".

base<sub>DS</sub>

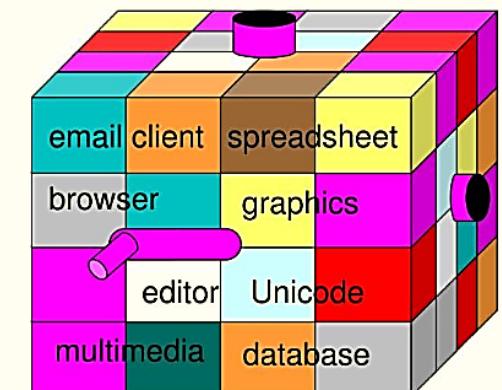
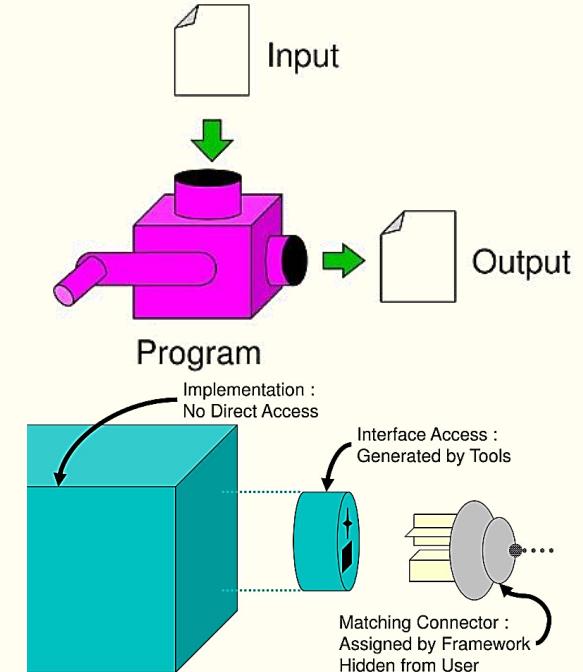
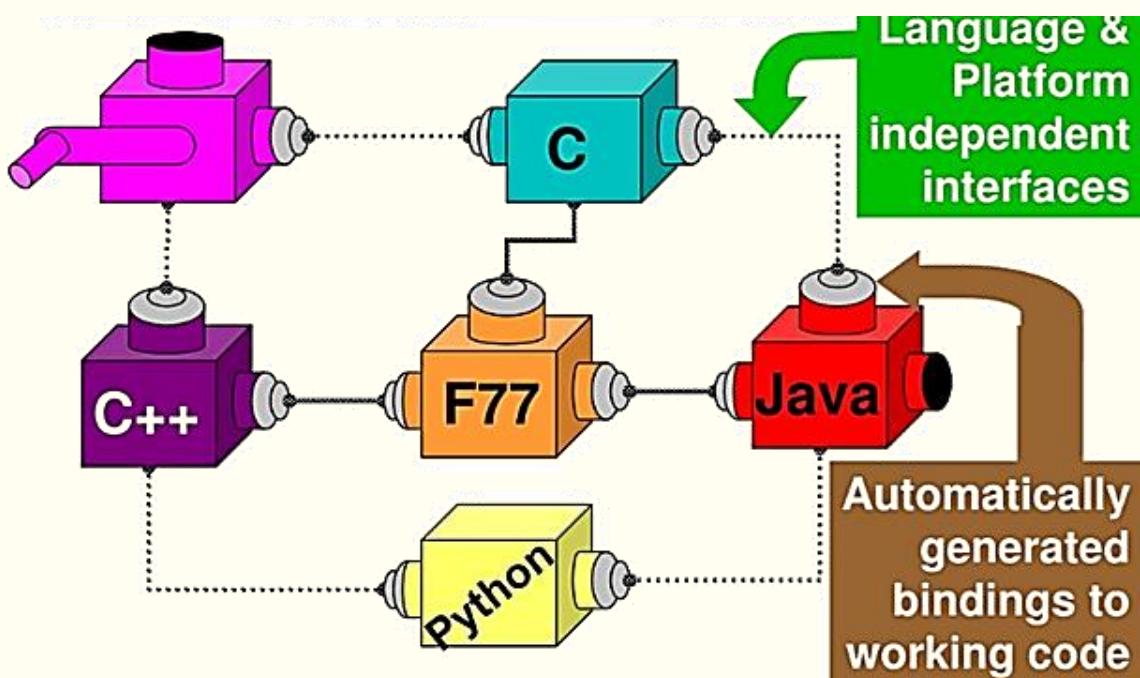
- The **degree** of Transparency:

- Engineers aim at **full** distribution transparency, **BUT**,
- There are communication **latency** issues and many covered **failures** either of **networks** or **nodes**.
  - You cannot distinguish a **slow** computer from a **failing** one,
  - You can never be sure that a server actually **performed** an operation before a **crash**,
  - This is a reality: engineers have no **global view** on the system.
- Full transparency will **cost performance** and will **expose system distribution**:
  - Keeping replicas **exactly** up-to-date with the master takes **time**,
  - For **fault-tolerance**, write operations need immediate flushing to storage disk, meant, data may vanish due to a crash.
    - For example, a customer makes a deposit to his bank account a few seconds before the bank system crashes.
- Distribution transparency is an **admirable** goal, but achieving it is **hard** work.



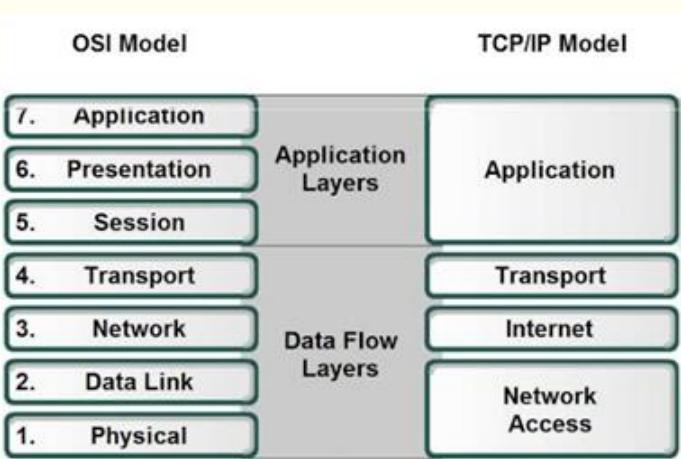
## • Openness of DS:

- Be able to **interact** with other **open systems**, irrespective of the underlying **software/hardware** components:
  - Systems should conform to well-defined **interfaces!!!**
  - **Discussion:** what are interfaces stereotypes?
  - Systems should easily **interoperate**.
  - Systems should support **portability** of applications.
  - Systems should be easily **extensible**.

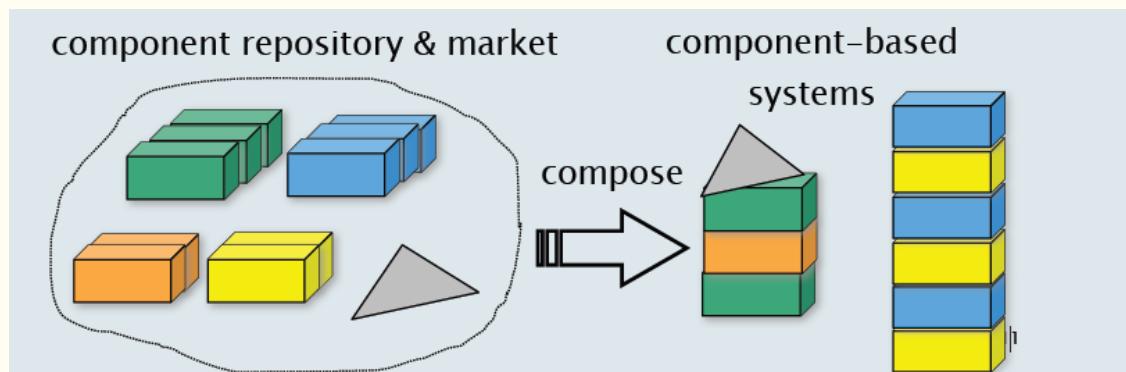


- Interoperability across multiple languages
- Interoperability across multiple platforms
- Incremental evolution of large legacy systems (esp. w/ multiple 3rd party software)

- Typically, system interoperability is implemented as **policies** and **mechanisms**.
- Engineers elicit for **DS** requirements using some questions such as:
  - What level of **consistency** do we require for client-cached data?
  - Which operations do we **allow** downloaded code to perform?
  - Which **QoS** to adjust for the varying bandwidth?
  - What level of security do we require for communication?

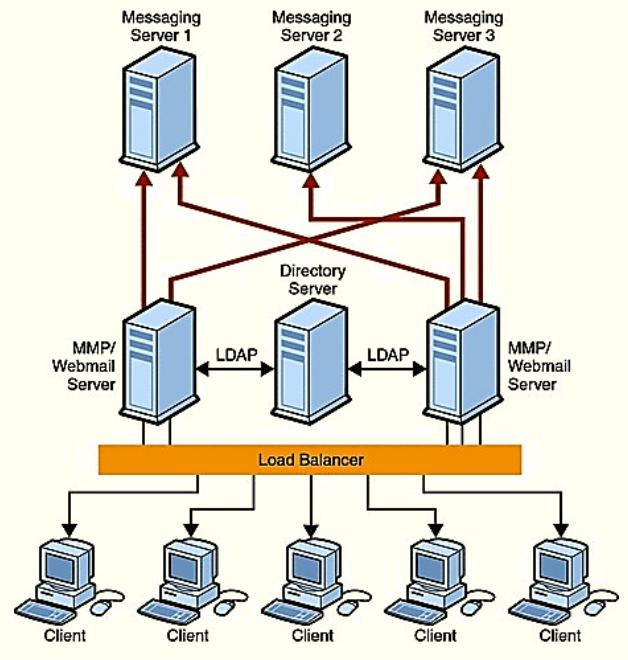
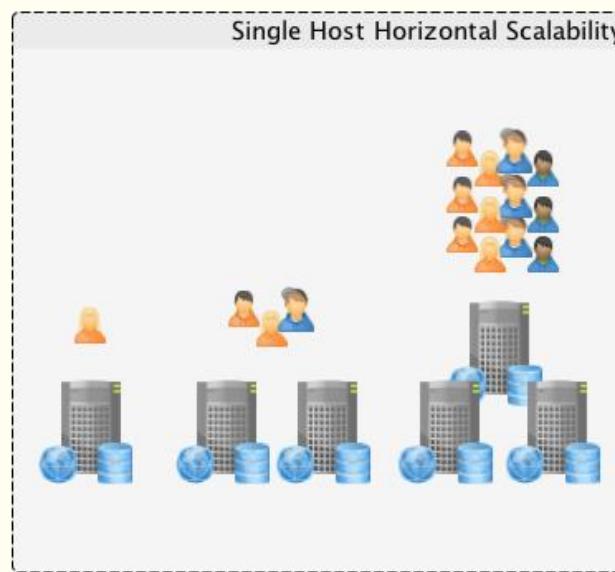
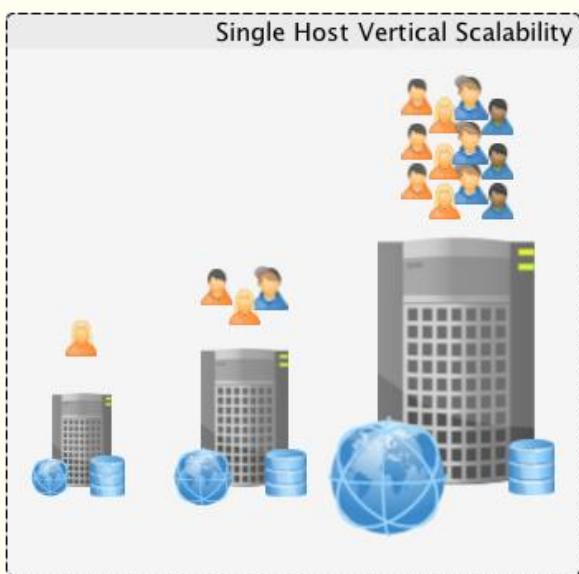


- The using of strict policies leads to many configuration parameters and complex management, thus, **balance openness** is needed.
- On recent approach is to build **component-based** systems.



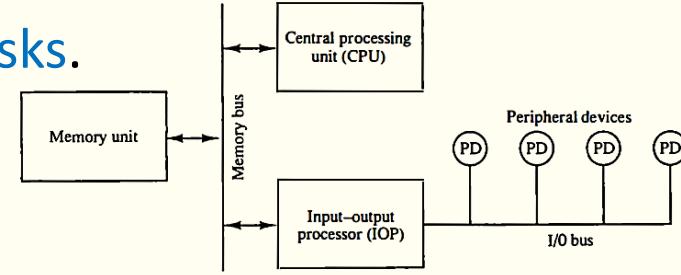
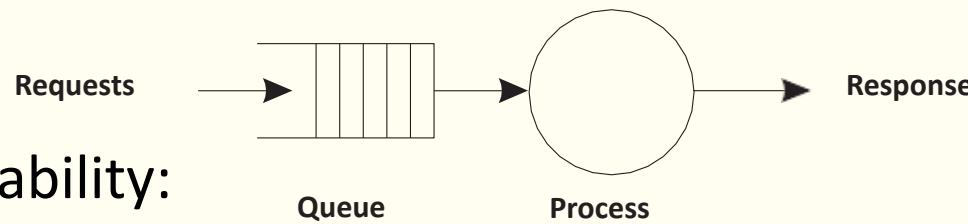
## • Scale of DS:

- Why DS actually **scales**?
- Often engineers refer to an increase in **size** by the term **scalability**, which is handled by **multiple powerful servers** operating **independently in parallel**. However, there are
  - An increase in the number of users and/or processes (**size scalability**),
  - An increase in distance between nodes (**geographical scalability**),
  - An increase in number of administrations (**administrative scalability**),
  - An increase in functionalities or loads (**load scalability**),
  - An increase in using components from different vendors (**heterogeneous scalability**).



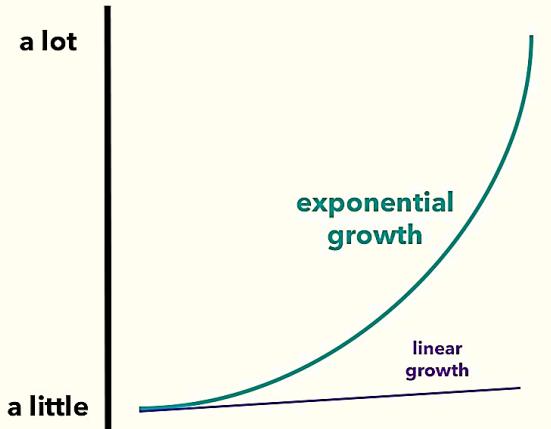
- Size scalability:

- Main issue for centralised systems,
- Computational capacity is limited by the CPUs,
- Storage speed is bounded by the transfer rate between CPUs and disks.
- Response time involves network delays and computational delays.



- Formal analysis for size scalability:

- How do engineers check whether to scale up their systems or not?
- A centralized service can be modelled as a simple queuing system:
  - The queue has infinite capacity,
  - $\lambda$  is the arrival rate of requests per time unit,
  - $\mu$  is the service rate of requests per time unit (processing capacity).
  - $\bar{N}$  is average number of request in the system  $\Rightarrow \bar{N} = \frac{\lambda}{\mu-\lambda}$ .
  - $S$  is average time to process  $\lambda$  requests in the system  $\Rightarrow S = \frac{1}{\mu-\lambda}$ .



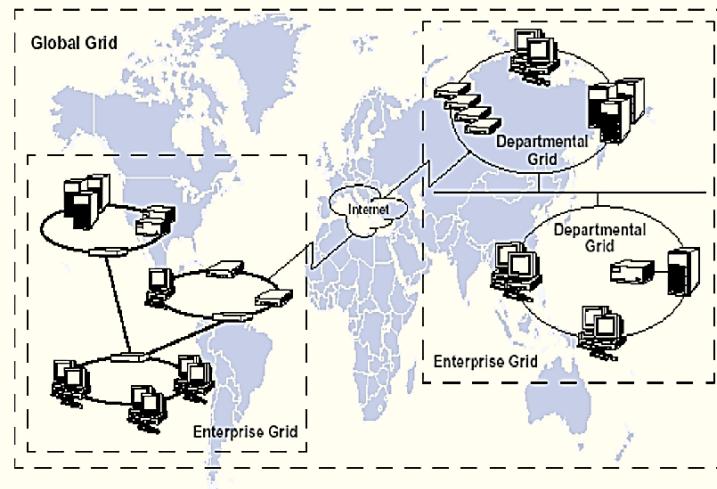
- As the number of requests increases linearly the service/response time increases exponentially.

- Geographical Scalability:

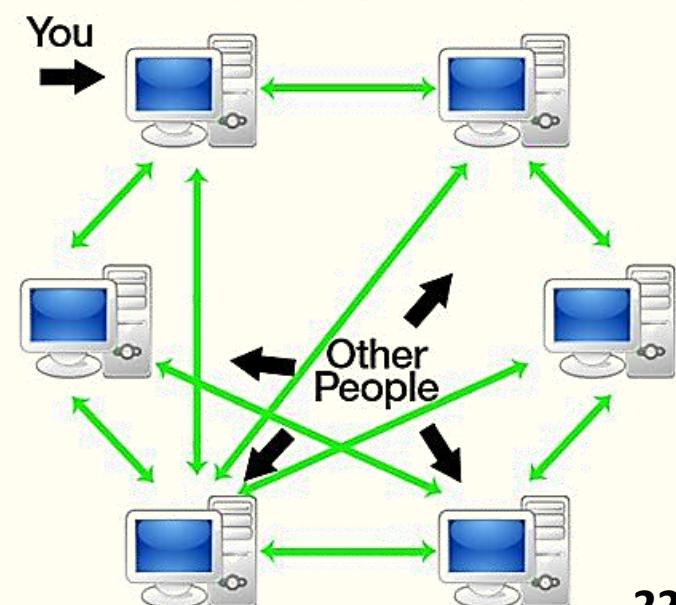
- Go from LANs to WANs is not simple, issues like synchronous client-server interactions may be prohibited due to network latency and reliability.
- Requires global naming services to support scalability and global broadcasting.
- Has special issues in CC (Edge & Grid).

- Administrative Scalability:

- A typical result of sharing expensive resources between different systems (usually centralised).
- Involves conflicting policies concerning usage, payment, management, and security.
  - Example: the control, manage, and use of a shared radio telescope constructed as large-scale sensor network.
- It, however, is not present in decentralised systems (Why?), i.e., peer-to-peer networks:
  - File-sharing systems (BitTorrent).
  - Telephony (Skype).
  - Assisted audio streaming (Spotify).
  - Cryptocurrencies (Bitcoins).



**Peer-to-Peer Model**



- Scalability Engineering:

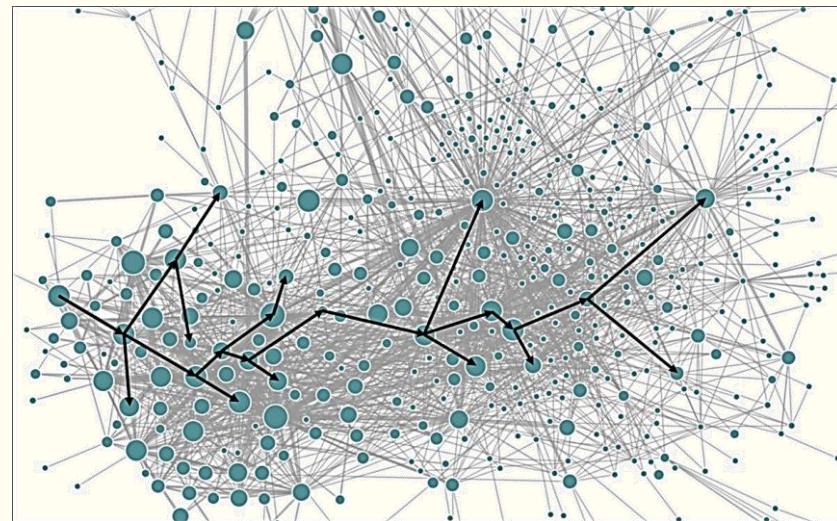
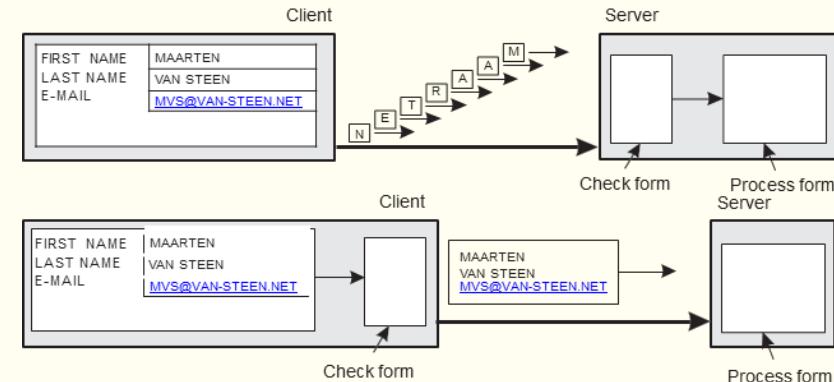
- Use **asynchronous** communication to overcome communication **latencies**.
- Use independent **services** for incoming **responses**.
- Move some **computations** to the **client** side.
- Implement **replication** and **caching**, e.g., file servers, databases, mirrored sites, and web caches.

- Limitations:

- One size doesn't fit all.
- Applying **replication** is problematic:
  - Inconsistent** copies, despite the frequent **updates**.
  - Requires **global agreement** and **synchronisation** on each modification.

- Notes:

- Global **synchronisation** blocks **large-scale** solutions.
- Tolerating consistency can reduce **global synchronisation** but this is an application dependent.

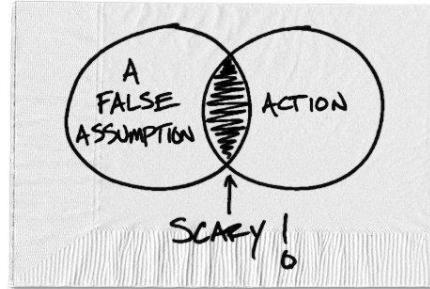
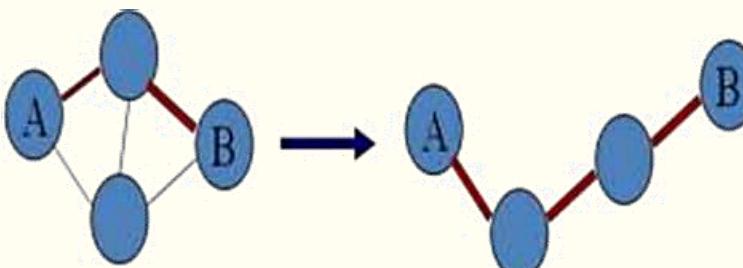


# Pitfalls in DS Engineering:

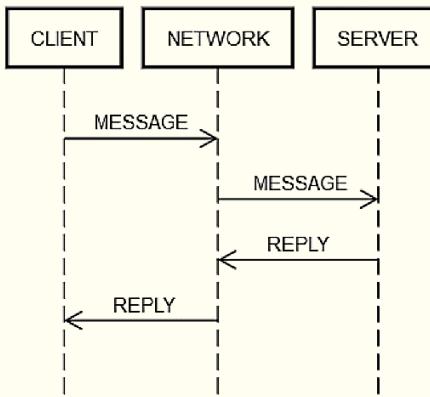
- Often, engineers make systems heavily complex and error-borne by applying **false assumptions**. Here is the Common **false** and **hidden** assumptions:

- The 8 fallacies, by L. Peter Deutsch, Sun Microsystems, 2015:

1. The network is **reliable**,
2. The network is **secure**,
3. The network is **homogeneous**,
4. The **topology** does not change,
5. **Latency** is zero,
6. **Transport** cost is zero,
7. **Bandwidth** is infinite.
8. There is one **administrator**.

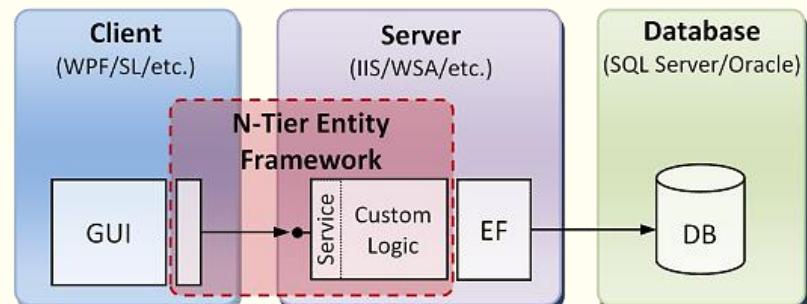
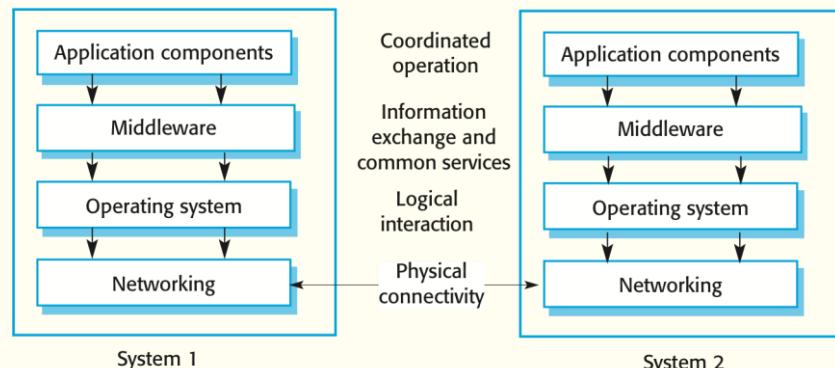
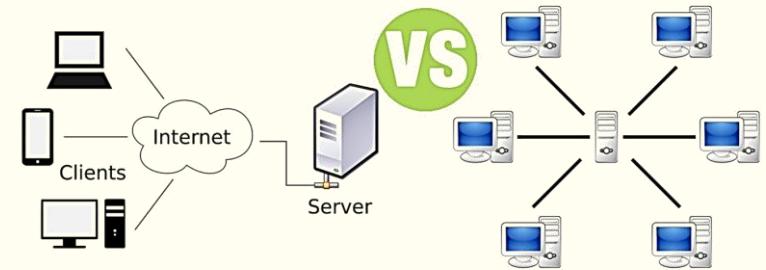


- Additional fallacies:
  - All **failures** have been considered in the development stage.
  - The server hardly **fails**, it is configured to **restart itself** no failure.
  - Engineers always have a **global view** of the distributed system.
  - A client can distinguish **network failure** from **server failure**.
  - All computing elements/nodes are **trustworthy**.
  - Using **formal analysis** and **modelling** is time-consuming and expensive, so, just skip it.



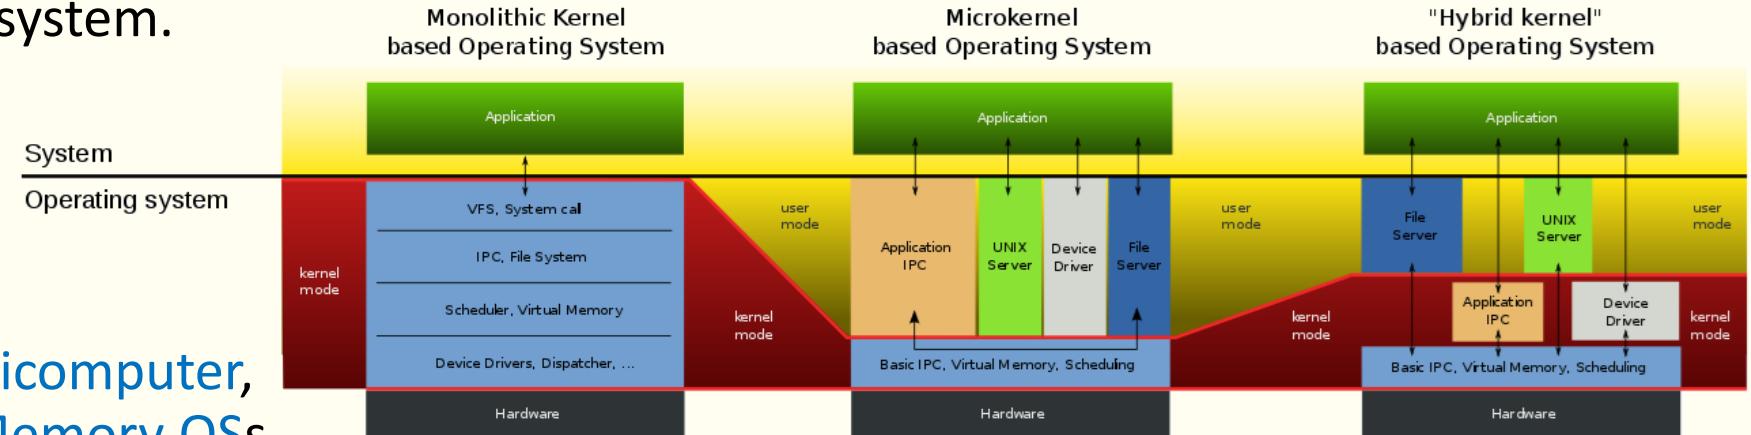
# Classification of DSs:

- Based on the architecture:
  - Client/Server systems (**Centralised**).
  - Peer-to-Peer systems (**Decentralised**).
  - Middleware systems.
  - Three-tier/N-tier systems.
- Based on the distribution of computing:
  - Operation Systems (**OSs**).
    - Distributed Operating Systems (**DOS**).
    - Network Operating Systems (**NOS**).
  - High-Performance Computing systems (**HPC**).
    - Cluster systems.
    - Grid systems.
    - Cloud systems.
  - Distributed information systems (**DIS**).
    - Distributed transaction processing.
    - Enterprise Application Integration (**EAI**).
  - Distributed pervasive systems.
    - Ubiquitous systems.
    - Mobile systems.
    - Sensor systems.



# • Distributed Operating System (DOS):

- DOS is a system software that comprises a collection of **independent**, **networked**, and physically **separate** computational **elements/nodes** ([Wikipedia](#)).
- Jobs/loads are serviced by multiple **CPU**s. Each node holds a specific software subset of the global aggregate operating system.

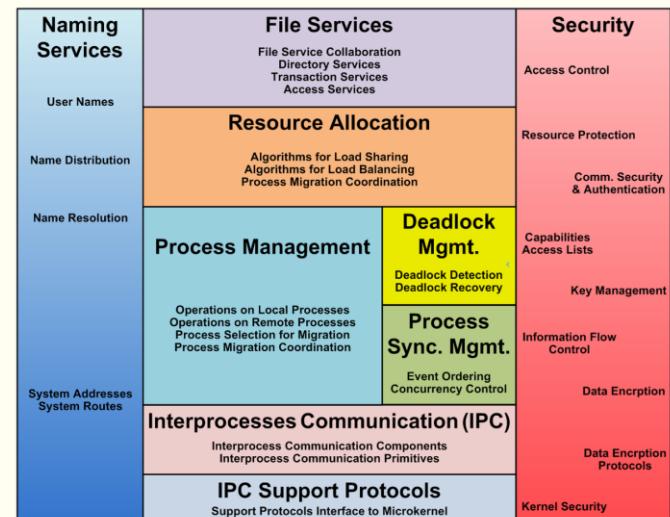


## • DOS Types:

Multiprocessor, Multicomputer,  
Distributed Shared Memory OSs.

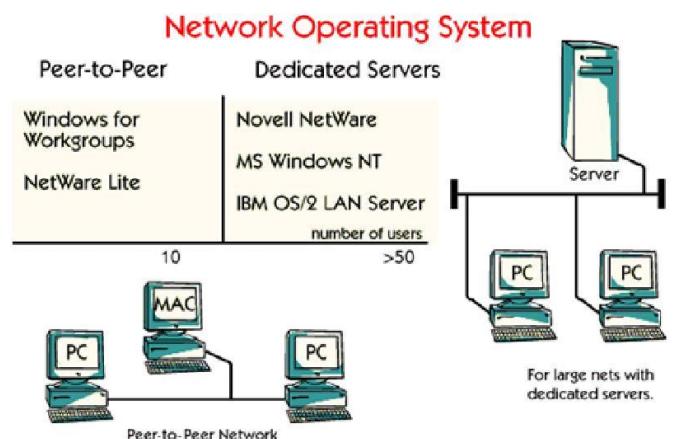
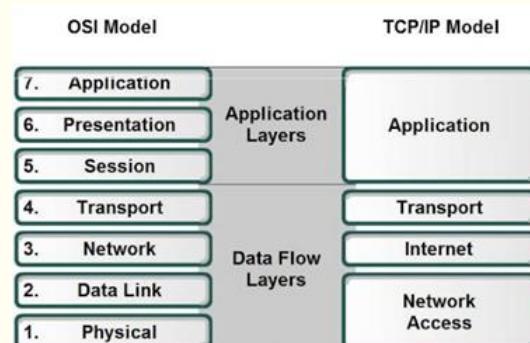
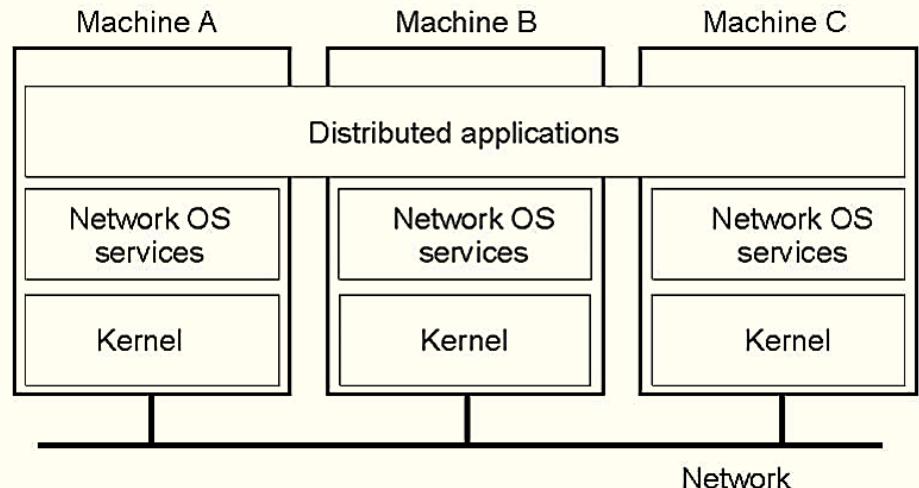
## • DOS Requirements:

- Performance & Scalability ([Intentional](#)).
- Distribution Management ([Complexity Issues](#)),
- Transparency: [Location](#), [Migration](#), [Replication](#), [Concurrency](#), [Scale](#),
- Openness ([Interoperability](#)),
- Reliability & Availability ([Fault-tolerance](#)),
- Synchronisation ([Atomicity](#), [Consistency](#), [Deadlock](#)),
- Flexibility ([Services](#), [Components](#)).



## • Network Operating System (NOS):

- NOS is a specialised OS for a network device such as a router, switch or firewall ([Wikipedia](#)).
- NOS enable computers nodes to participate and share files and printers access within the network.
- Proprietary & Open Source.
- Client/Server & Peer-to-Peer.
- NOS Requirements:
  - Resource Sharing & Cost Saving ([Intentional](#)).
  - Transparency: [Location](#), [Platform](#).
  - Openness ([Interoperability](#)),
  - Reliability & Availability ([Fault-tolerance](#)),
  - Synchronisation ([Consistency](#)).
  - Security.



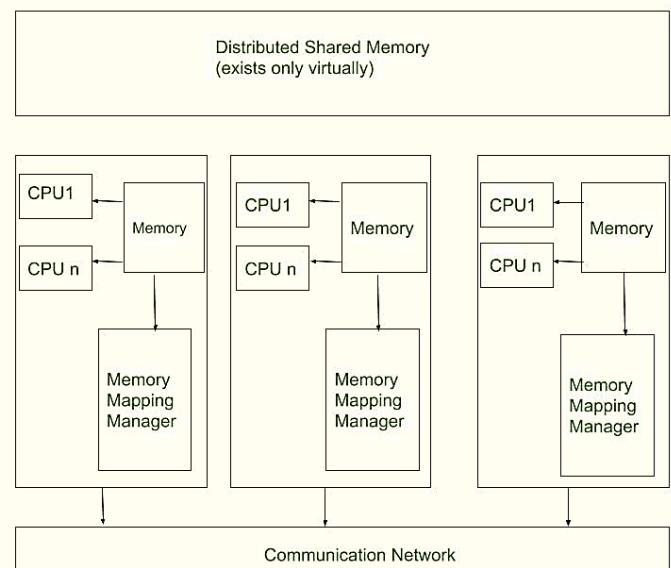
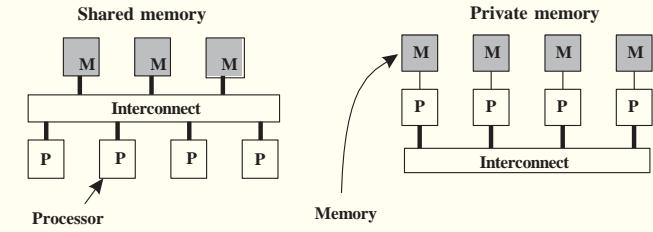
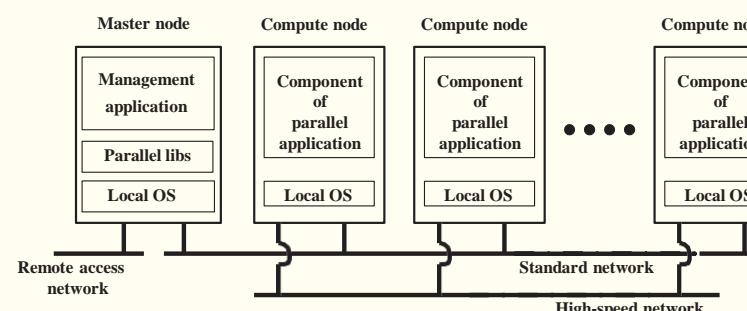
## • Limitations:

- DOS, complex & non-flexible,
- NOS, not a coherent computer,

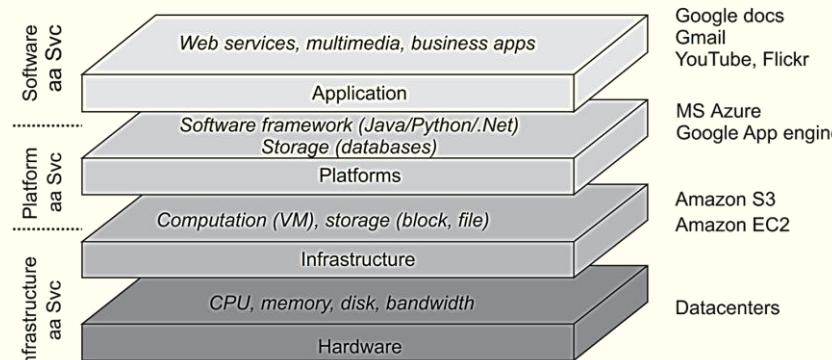
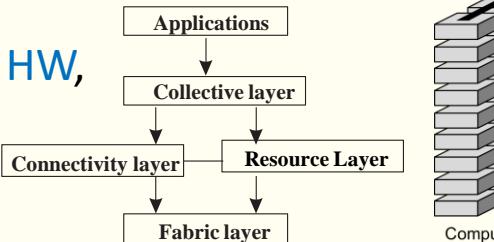
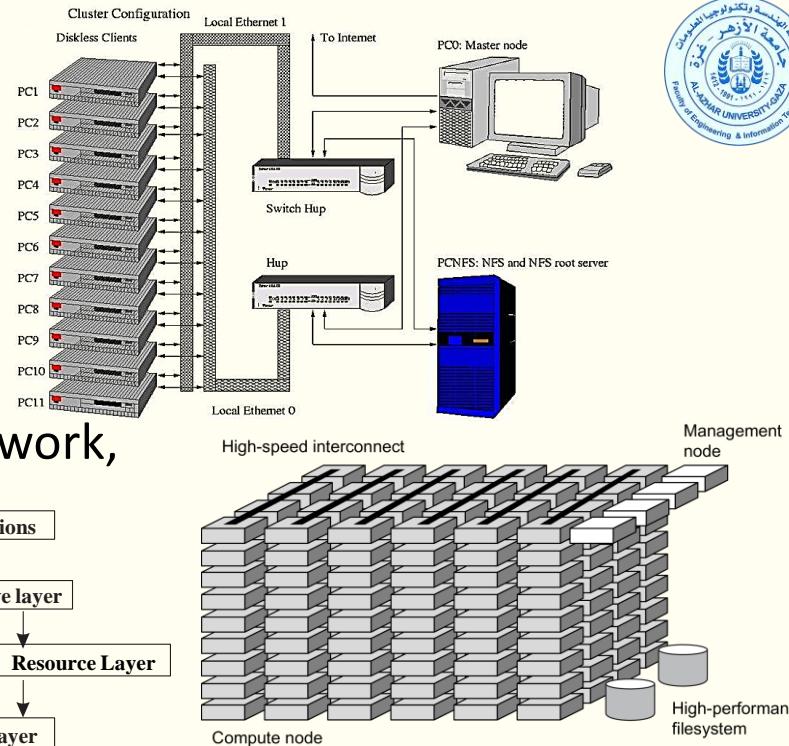
• For the best of both, consider the **Middleware** approach.

## • High-performance distributed computing (HPC):

- Typically is characterised by parallel computing, multiprocessor/multicore, and high-end multicompiler.
- Typically, multiple processes/threads run at the same time;
- All threads have access to shared data; access is controlled via synchronisation mechanisms (e.g., Semaphores).
- However, this approach does not easily scale due to hardware limitation (i.e., Mainboards have limited sockets for CPUs & RAMs) .
- A solution was to use shared-memory model on top of multicompiler.
- The shared-memory model suffers from synchronisation (i.e., concurrent access) and message-passing (i.e., latency) issues.
- Later, the Distributed Shared-Memory (DSM) model is introduced.  
In the DSM, a processor can address a memory location as if it were local memory (Virtual-Memory Technique). However, DSM has been abandoned due to performance issues.
- Implementations of the HPC:
  - Cluster computing.
  - Grid computing.
  - Cloud computing.



- **Cluster computing:**
  - A group of high-end **computers/servers** connected via a **LAN**,
  - **Homogeneous**, nodes use the same **OS** (i.e., **DOS**) and have identical hardware, There are a **single** managing **node**.
- **Grid computing:**
  - **Heterogeneous**, lots of nodes of different OSs and hardware,
  - Dispersed across several organizations and may span a **wide-area network**,
  - Architecture for grid computing:
    - **Fabric** layer provides interfaces to local resources, i.e., **OS & HW**,
    - **Connectivity** layer runs communication protocols,
    - **Resource** layer manages a single resource,
    - **Collective** layer handles multiple resources, e.g., discovery, scheduling, & replication,
    - **Application** layer is actual grid applications in a single organisation.
- **Cloud computing:**
  - General architecture:
    - **Hardware** layer includes processors, routers, and power **HW**.
    - **Infrastructure** layer deploys **virtualisation** techniques.
    - **Platform** layer provides higher-level abstractions for storage & so (i.e., **Middleware**).
    - **Application** layer.
- **Discussion:** what engineering requirements do **HPC** systems need?



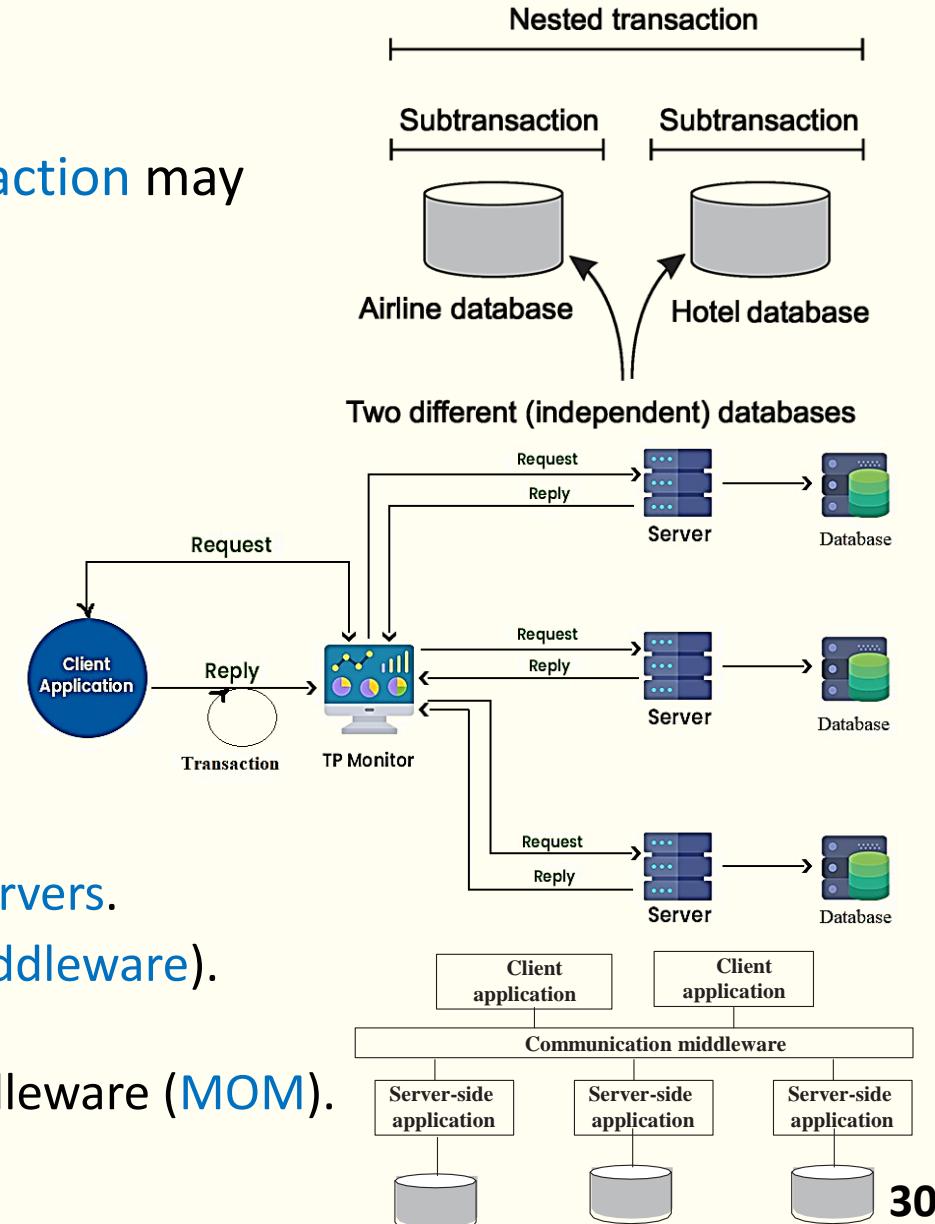
- Distributed information systems (DIS):

- At present, organisations management **relies** on many networked applications, e.g., DBMS, CRM, HRM, FMIS, ... etc.
- DIS** diversity makes achieving **interoperability** painful!!!
- But, **Integration** in the DIS is needed because a **single transaction** may execute over **many applications/machines**.
- Distributed transaction processing:

- Transaction Issue: **all-or-nothing**,
- Atomic**: happens indivisibly,
- Consistent**: does not violate system invariants,
- Isolated**: not mutual interference,
- Durable**: **commit** means changes are permanent,
- TPM**: Transaction Processing Monitor.

- Enterprise Application Integration (EAI):

- Data involved in a **transaction** is **distributed** across several **servers**.
- A **TP Monitor** coordinates the execution of a transaction (**Middleware**).
- Middleware** offers facilities for integration:  
i.e., Remote Procedure Call (**RPC**)! & Message Oriented Middleware (**MOM**).



- **Distributed pervasive systems:**

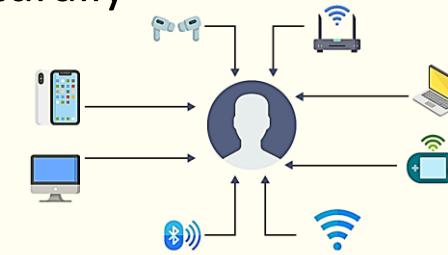
- Systems in which nodes are **small, mobile, embedded** in a larger system, and naturally **blend** into the user's **environment**.
- The distinction is the usage of **wireless networks, sensors , actuators**, and **IoT**.
- **Ubiquitous computing systems:**
  - **Distribution:** devices are networked, distributed, and accessible transparently.
  - **Interaction:** interaction between users and devices is **appreciable & continuous**.
  - **Context awareness:** system is aware of a user's context.
  - **Autonomy:** devices operate **autonomously** and are **self-managed**.
  - **Intelligence:** system handles a wide range of **dynamic actions** and **interactions**.

- **Mobile computing systems:**

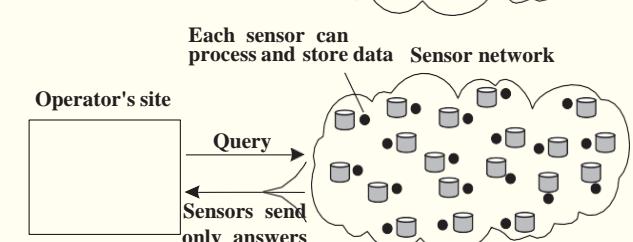
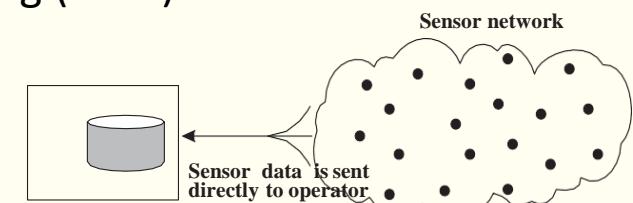
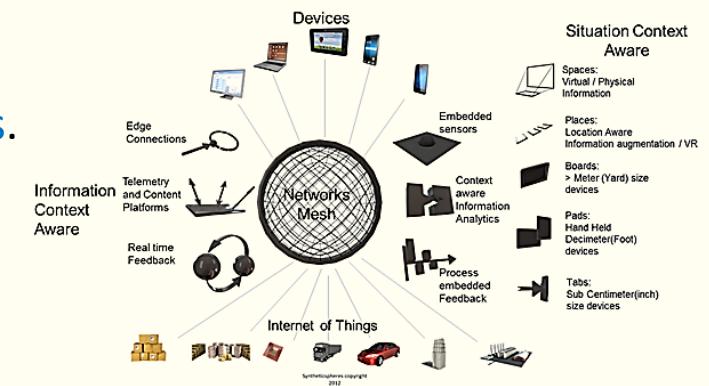
- Consider systems for smartphones, tablets, GPS devices ... etc.
- They are **mobile, location changes** over time, thus **services change** too,
- Got issues like **connectivity, reachability, and discoverability**... etc.
- Known architecture: Mobile Edge Computing (**MEC**) and Mobile Cloud Computing (**MCC**).

- **Sensor computing systems:**

- A collaborative system that **senses** and **actuators**.
- Consists of many (**10s-1000s**) simple nodes.
- Data aggregation either **centralised** or **distributed**.
- Known issues like **power, failure, connectivity**.
- Decentralised dissemination techniques are used for **fault-tolerance**.



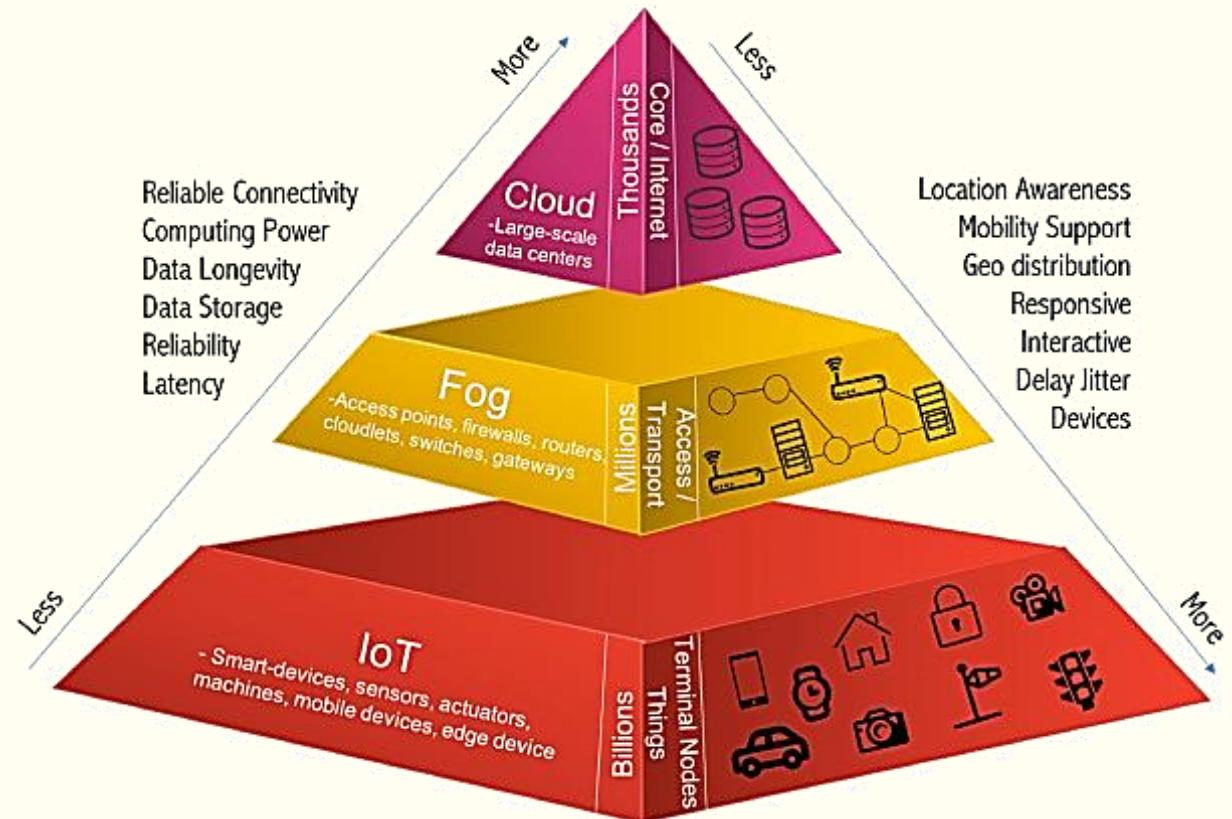
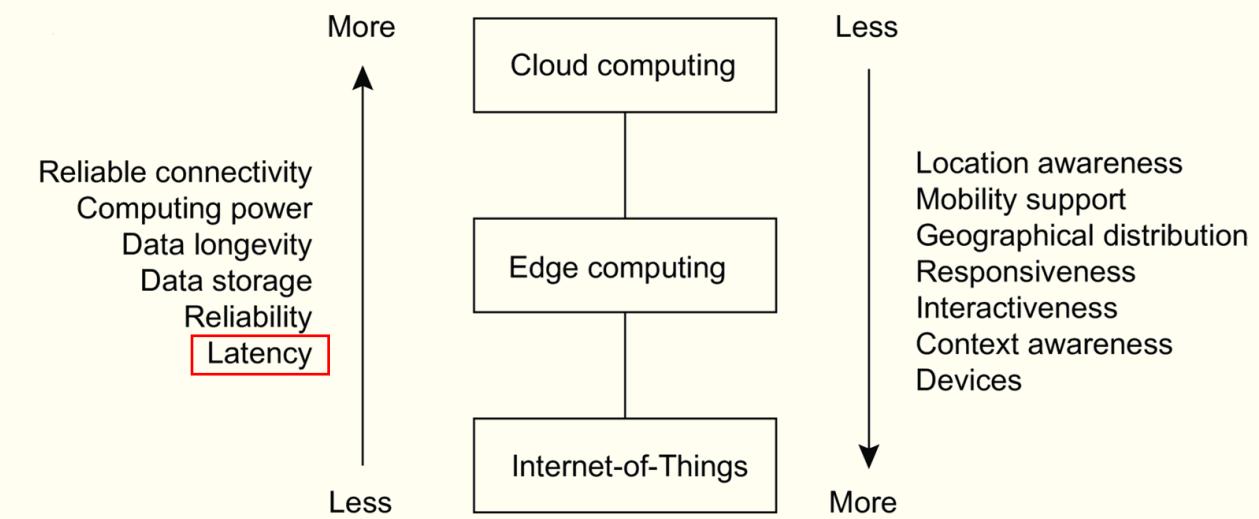
**Ubiquitous Systems**



- **Discussion:** what requirements does pervasive systems engineering need?

- A hierarchical view from clouds to devices:

- Clearly, DSs span a wide diverse range of networked computer systems, typically connected through a LAN or recently WAN (i.e., IoT).
- Higher up the hierarchy, the qualities improve with a less number of computers;
- Lower in the hierarchy, there are location-related and latency issues. Also, the number of devices increases.

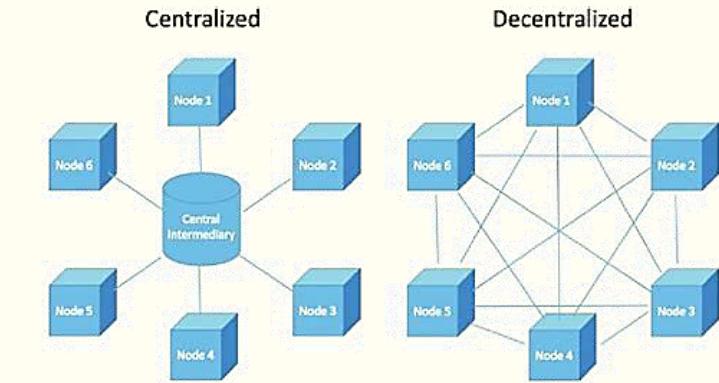


Yousefpour A., et al. All One Needs to Know About Fog Computing and Related Edge Computing Paradigms: A Complete Survey.

# DISTRIBUTED SYSTEMS ARCHITECTURES

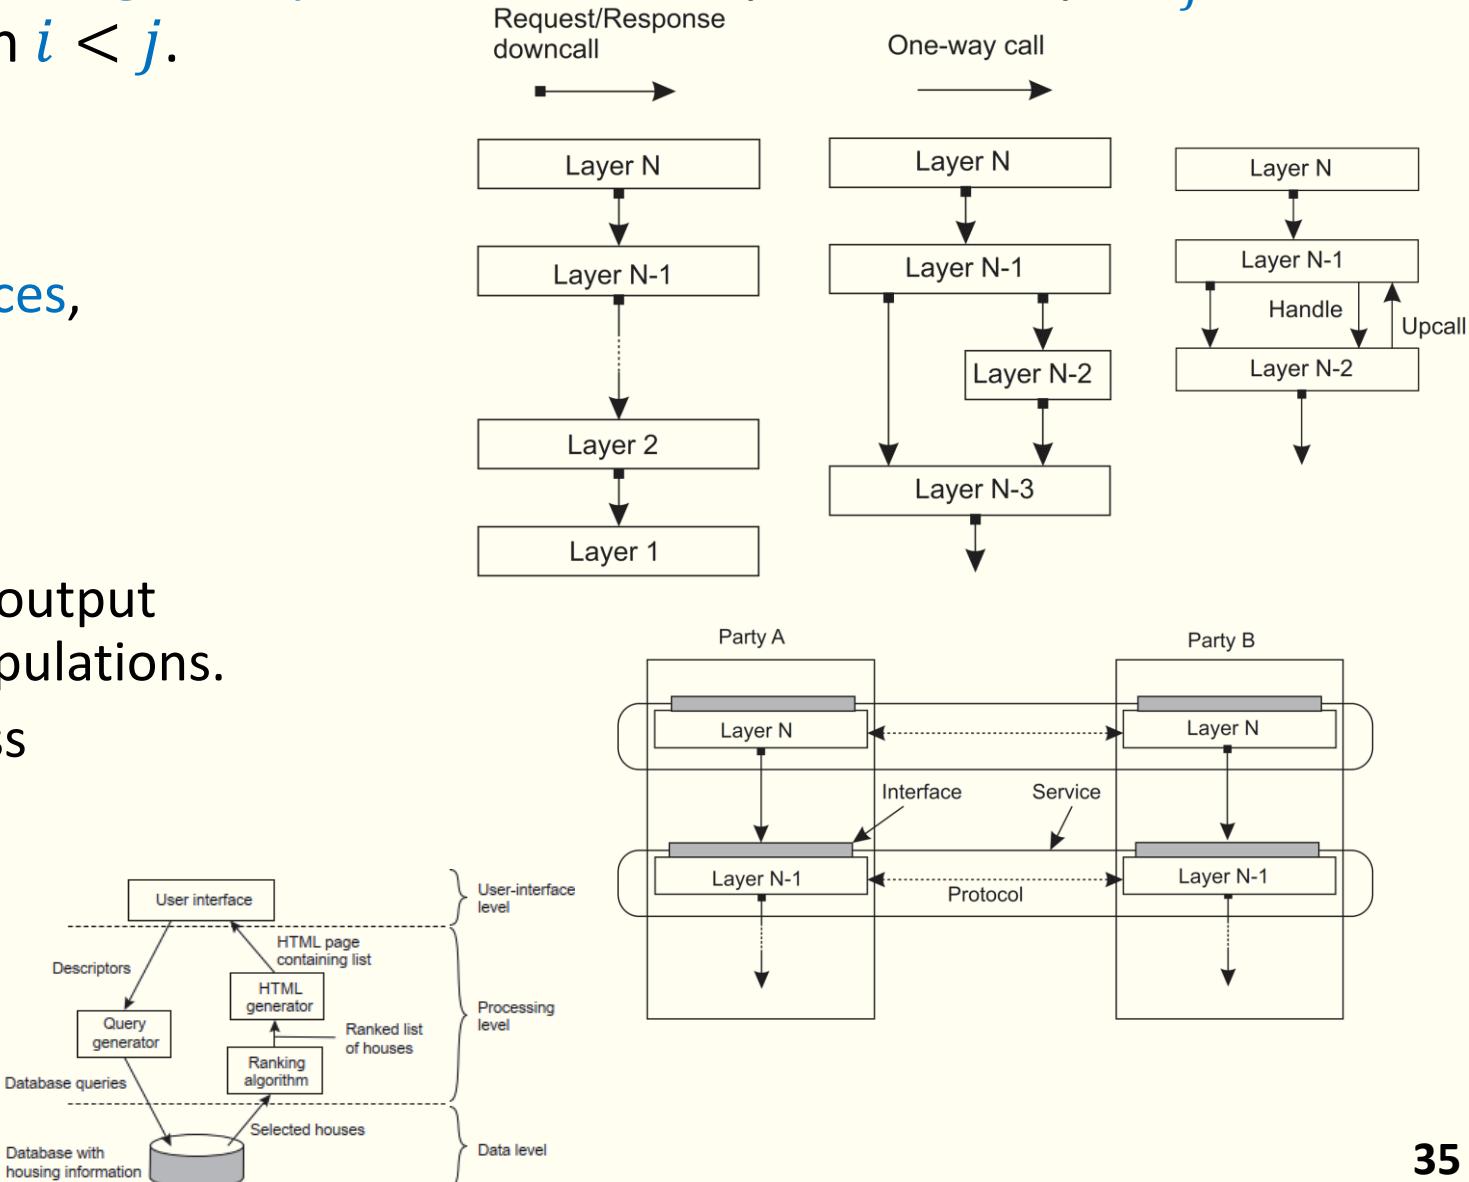
# Architectural styles:

- A DS is built of interconnected computing units (i.e. services & components). Units are meant to be replaceable.
- A DS style is formulated in terms of:
  - components with well-defined interface,
  - the way that components are connected to each other, (i.e. Connectors),
  - the data exchanged between components,
  - how these components and connectors are jointly configured into a system.
- Styles:
  - Layered & Multitiered Architectures.
  - Service-oriented Architectures (SOAs).
  - Publish-subscribe Architectures.
  - Middleware Architectures.
  - Symmetrically DSs Architectures.
  - Hybrid System Architectures.
- Discussion: Is adopting the right architecture crucial for the success of large software development?



# Layered Architectures:

- Simply, systems **units** are **organised** into **logical layers** where a component at layer  $L_j$  can make a down call to a component at  $L_i$ , with  $i < j$ .
- Layered communication protocols:
  - Connect parties via protocol stack,
  - Characterised by **interfaces** and **services**,
  - e.g. **TCP/IP**, **HTTP**, **Web systems** ...
- Application layering:
  - The **Interface Layer** consists of input/output elements, and elementary data manipulations.
  - The **Processing Layer** handles business logic and major data manipulations.
  - The **Data Layer** is the data storage.
- Drawbacks:
  - Strong **dependency** among layers.

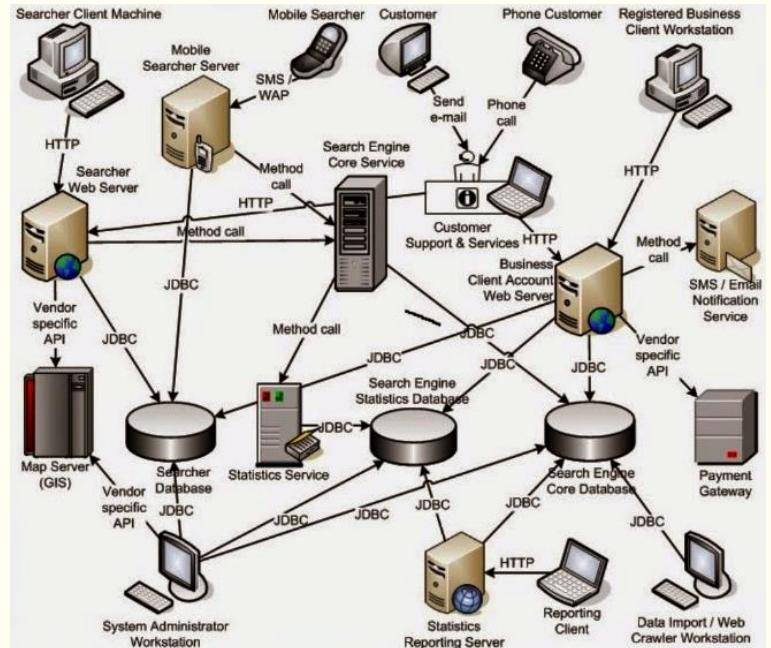
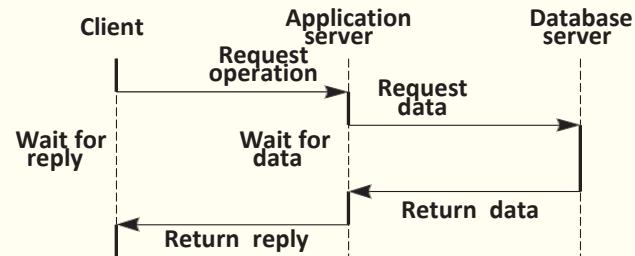
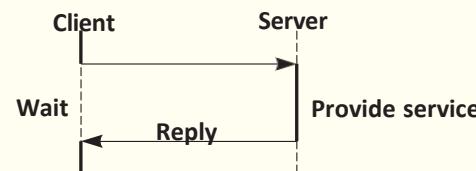


# Layered-system (Multitiered) Architectures:

- This style is classified by considering where software components are located.

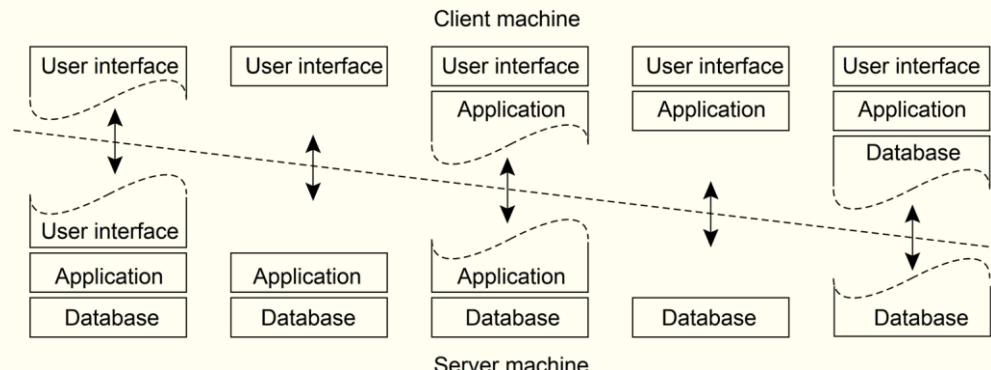
## Client/Server Architecture:

- There are processes offering services (**servers**) and processes that use services (**clients**),
- Interaction follows **request/reply** model with a **connectionless** or **connection-oriented** communication.
- Distinction between **client** and **server** is problematic (**Why?**).



## Multi-tier Architecture:

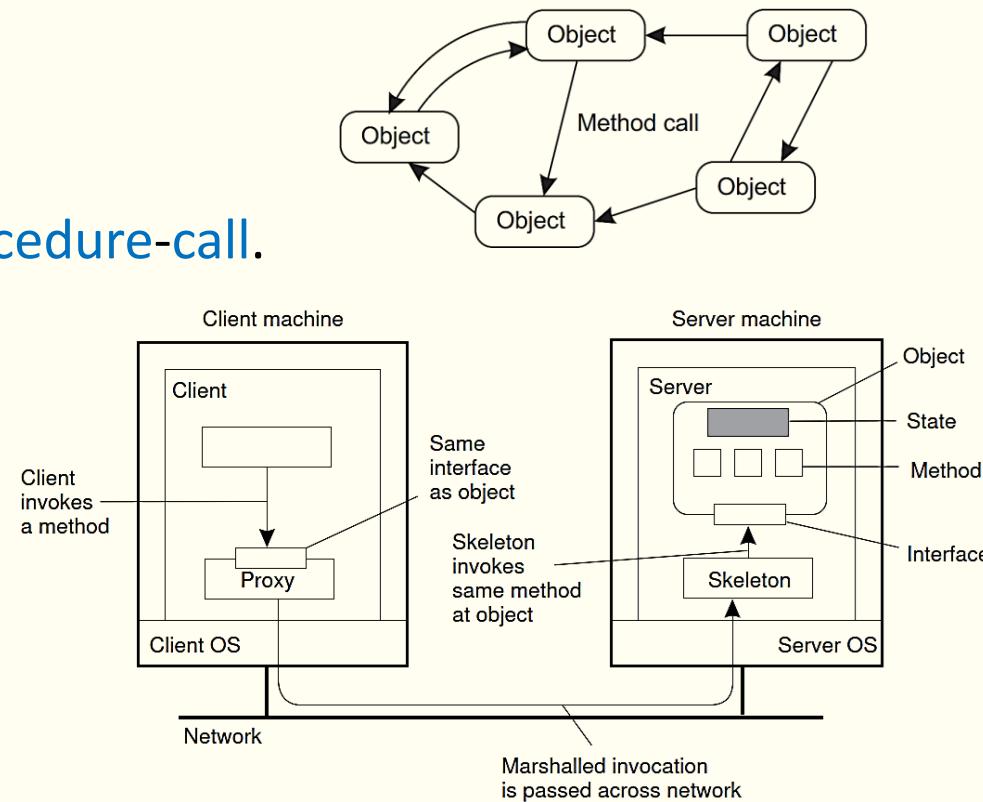
- Tiers number impacts performance.
- It can have various organisations.
- Example: Computing Terminals, **NFS**, Web.



- Discussion :** what factors to consider in choosing a multi-tier style?

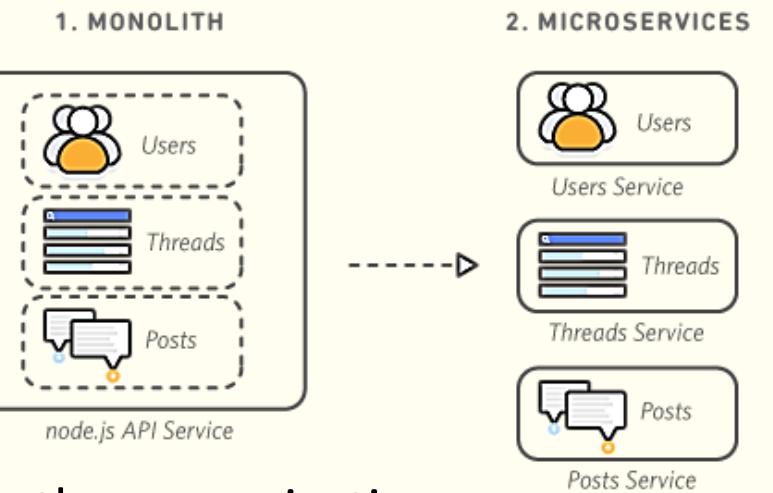
# Service-oriented Architectures (SOAs):

- More loose organisation into a collection of separate, independent entities (**Objects, Components, Services**).
- Object-based style:
  - Distributed Independent **objects** connected through a **procedure-call**.
  - Objects **encapsulates** data and offer **methods** without revealing the internal **implementation**.
  - At **client-side**, to bind a remote object:
    - Load **Proxy (Stub)** object to memory,
    - Proxy object **packs** remote **invocations (Requests)**.
  - At **server-side**, a **skeleton** object is always up,
    - Skeleton object **unpack** requests, **pass** them to server objects, **packs** and **sends** responses.



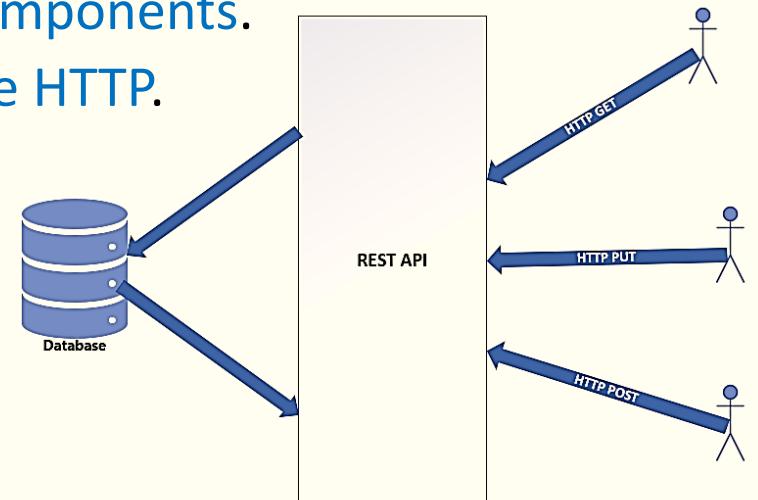
- Microservice style:

- In a **service-oriented** architecture, a **DS** is constructed as a composition of many **services**.
- A service is a **self-contained** entity.
- Provides a specific **functionality** (service).
- Can be **queried** and **access** remotely, but has **complexity** and **integration** issues.
- e.g., e-shops: interface, orders, & database, Also, **payment & delivery**; which typically are services belonging to other organisations.



- Resource-based style:

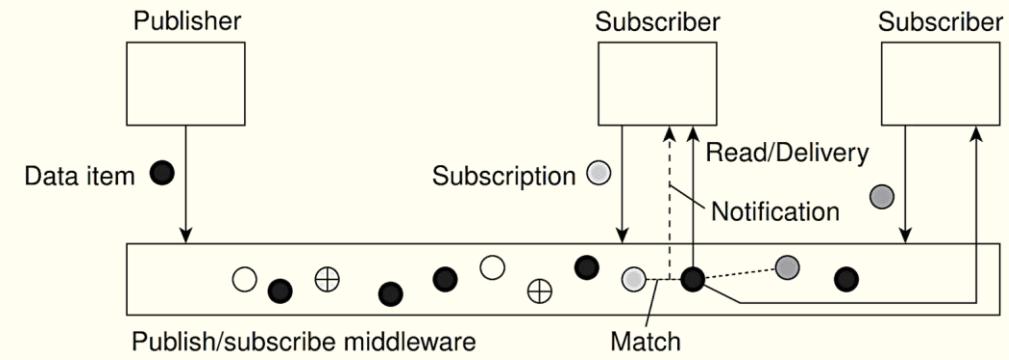
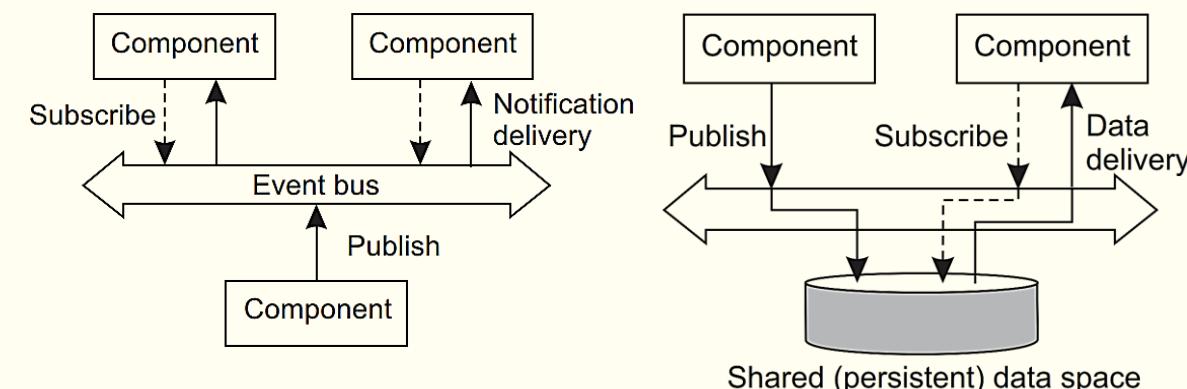
- A **collection of resources** that are individually managed by **components**.
- **Stateless, fire-and-forget client/server** entities with **cacheable HTTP**.
- It is widely used for the **Web** and is known as **Representational State Transfer (REST)**.
- **Single naming scheme, uniform interfaces**, and **self-described** messages.
- **Simple, but stateless**.
- Example: sharing a cloud file using a **URI link**.



# Publish-Subscribe Architectures:

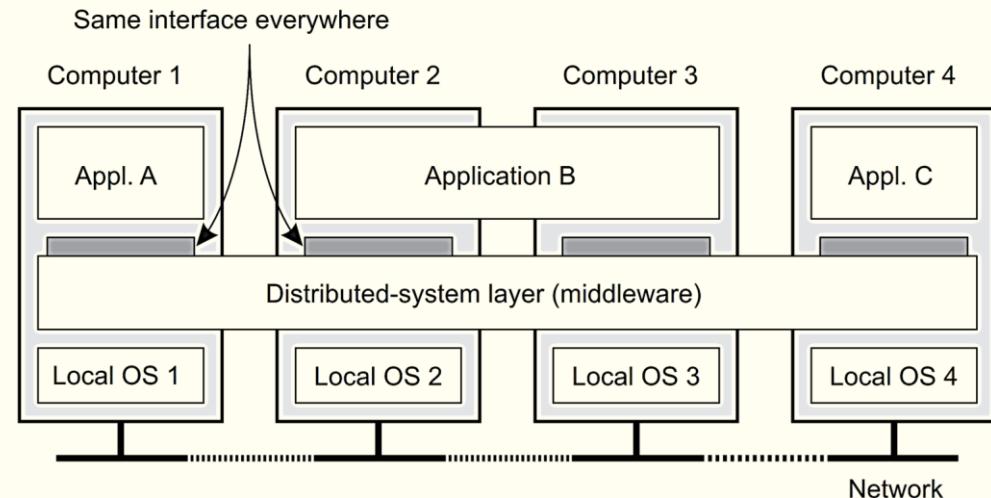
- Some large DS demand architecture with **dependencies as loose as possible**.
- So, a strong **separation between processing and coordination** is needed, this has led to:
  - **Direct coordination**, when processes are **temporally and referentially coupled**,
  - **Event-based coordination**, when processes are **temporally or/and referentially decoupled**.
  - Processes do not know each other explicitly! So, processes communicate via **publish-notification** fashion.
  - Processes may need to **subscribe** to receive notifications.
- Also, there is **content-based publish-subscribe** systems.

	Temporally coupled	Temporally decoupled
Referentially coupled	Direct	Mailbox
Referentially decoupled	Event-based	Shared data space

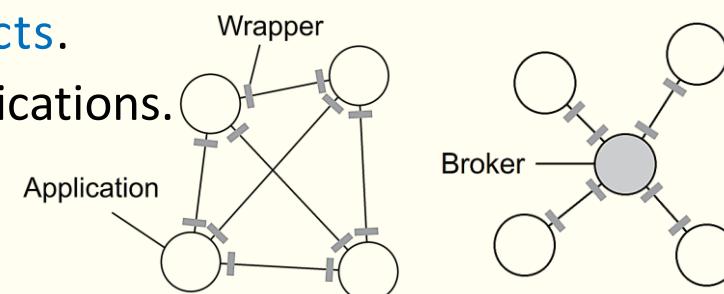


# Middleware Architectures:

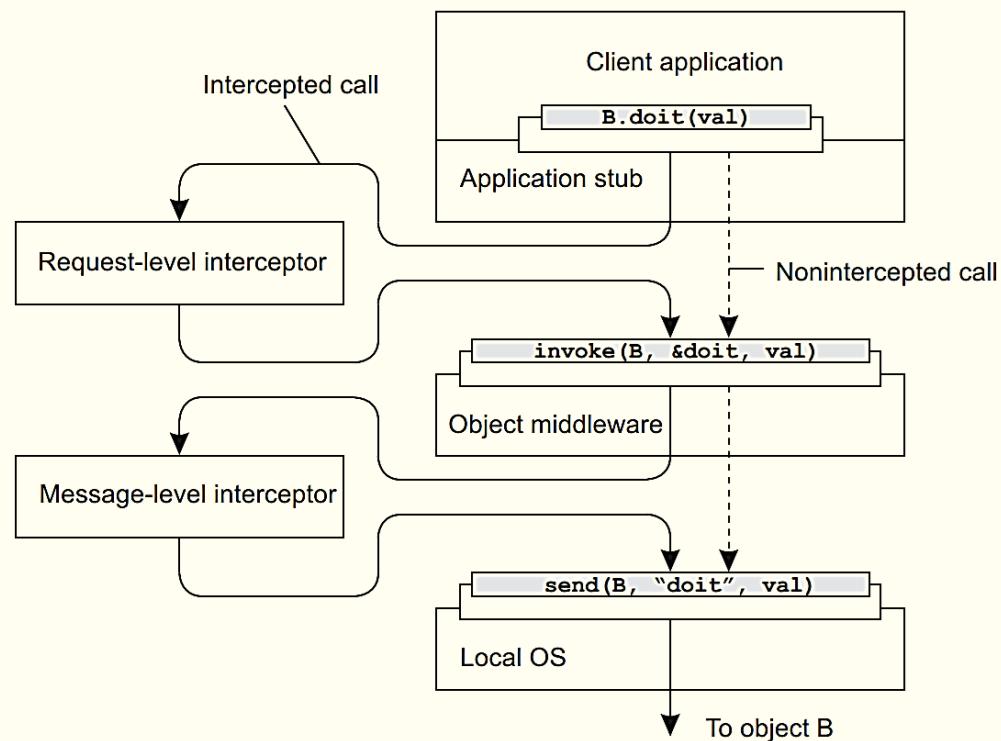
- In 1990s, to ease the development of **DS**, a separate layer of software is logically placed on top of the respective **OS/platform** known as **Middleware**.
- Middleware facilities inter-communication, security, accounting, masking **OSs**, platforms, failures, and recovery.
- Middleware must be **portable, scalable, interoperable, transparent, and fault-tolerant**.
  - **Discussion:** **OSs** are a sort of mindware, aren't they?



- Middleware construction involves one or more:
  - **Wrapper** is an interface **component/object** to a certain set of client applications.
  - **Adapter** is a **component** that allows applications to invoke **remote objects**.
  - **Broker** is a **centralized component** that handles all the **accesses** among applications.
  - **Interceptor** is a **component** that break the flow-of-control to adapt **middleware** to a need of an **application**.

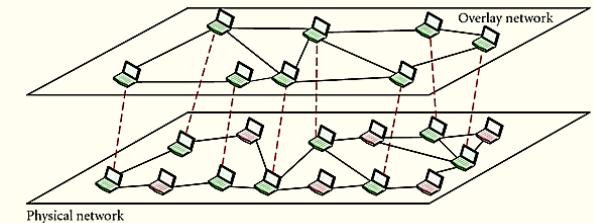
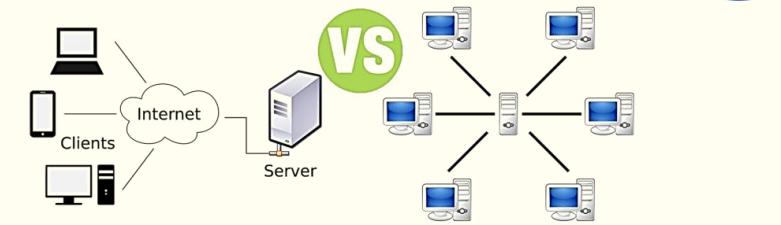


- Modifiable Middleware:
  - Computing environment **evolves rabidly**, e.g., **Mobile** and **Pervasive** systems.
  - Advanced middleware should **handle** and **hide context changes** and **events** on-the-fly.
  - Organizing middleware to be modifiable requires more research work.
- **Component-based design** focuses on supporting modifiability through composition. A system may either be configured **statically** at **design time**, or **dynamically** at **runtime**.



# Symmetrically DSs Architectures (Decentralisation):

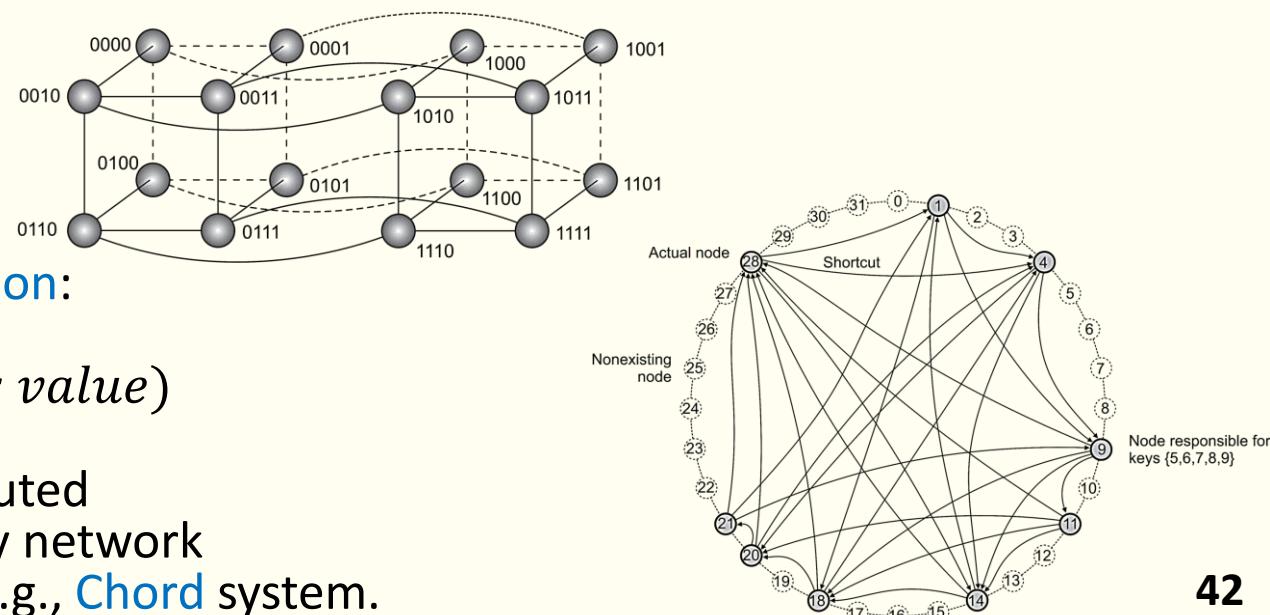
- The distribution of computation in DSs takes two forms:
  - Vertical** distribution refers to splitting the computations into **Master-Slave** models corresponding to the application organisation i.e., **client/server** and **multitiered** architectures.
  - Horizontal** distribution refers to **equivalent** computations models of **client/server** architecture corresponding with application organisation, e.g., **peer-to-peer** systems.
    - The interaction between processes/nodes is **symmetric**, “each process/node will act as a **client** and a **server** (a.k.a **servant**) at the same time.
    - Due to the **symmetric** style, **peer-to-peer** architectures revolve around how to organize the processes/nodes in an **overlay network**.



- Structured P2P systems:**
  - Processes/nodes are organised in an **overlay** that adheres to a **deterministic** topology.
  - Make use of a **semantic-free** index to uniquely **identifies** nodes with a **key** and a **global hash function**:

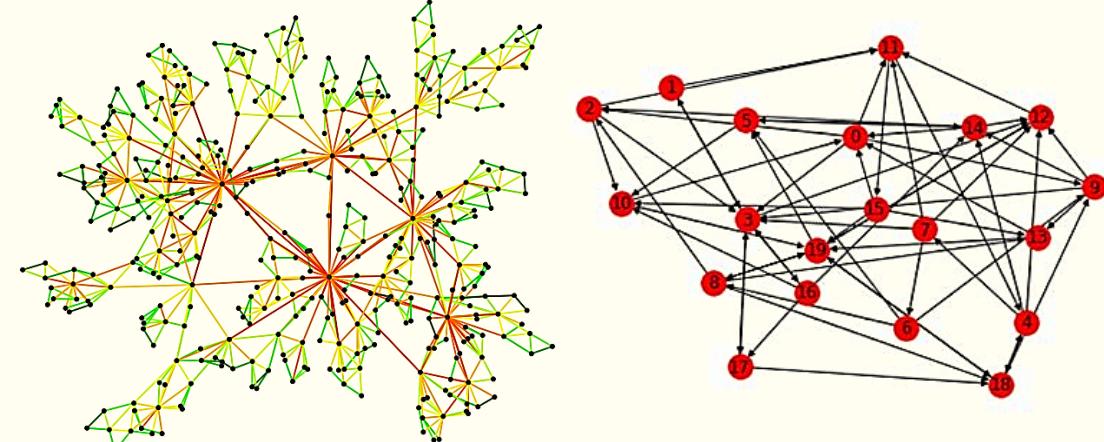
$$\text{key}(\text{data item}) = \text{hash}(\text{data item's value})$$

- The Distributed Hash Table (DHT) is vital for distributed computation, data access and lookup in the overlay network (**IP** might be present in the underlying network!), e.g., **Chord** system.

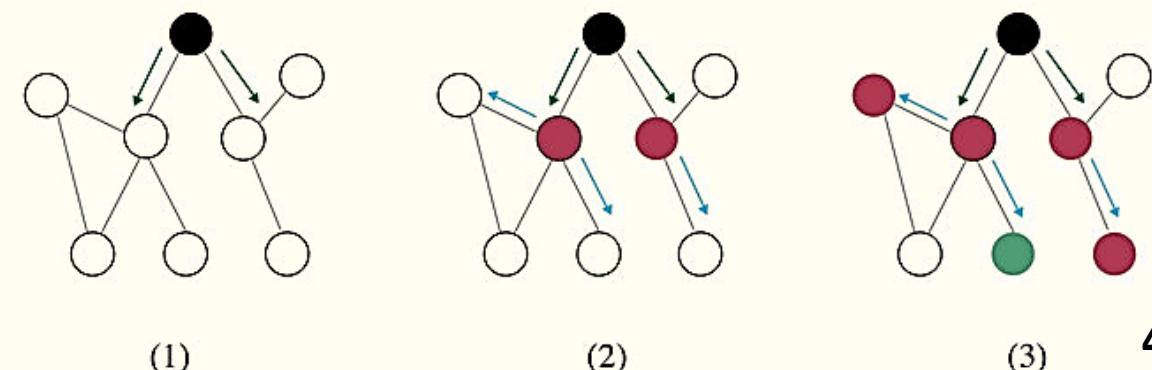


- Unstructured P2P systems:

- Each process/node maintains an ad-hoc/stochastic list of neighbours.
  - The overlay network forms a random graph with an edge  $(u, v)$  exists only with a certain probability  $\mathbb{P}[\langle u, v \rangle]$ .
  - P2P overlay can be static but commonly it forms dynamic topologies.
  - Exhibit several issues:  
node joining, neighbours membership, and search/lookup, routing, ... etc.

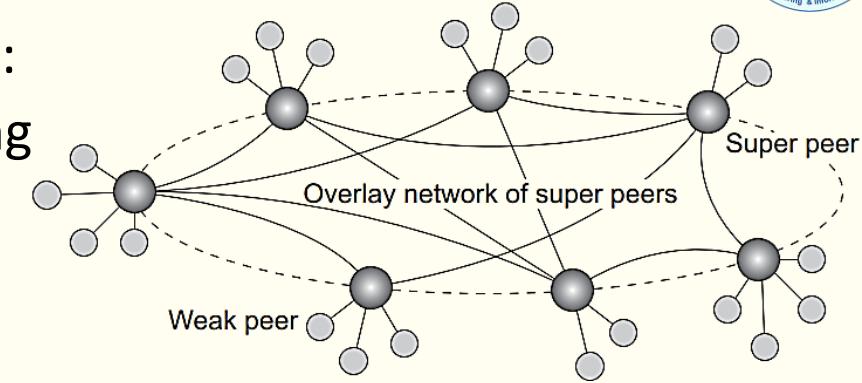


- Diffusion/Dissemination techniques in P2P overlays:
    - Flooding: node  $u$  requests all neighbours, node  $v$  ignores if seen, forwards if not applicable, responds directly to  $u$  or to precedence node.  
To avoid expensive flooding, limit search by Time-To-Live (TTL) value.
    - Random walk: node  $u$  requests a random neighbour node, forwards randomly if not applicable, or responds directly to  $u$ .  
To decrease lookup time, issuer node starts  $n$  random walks simultaneously.  
To stop random walk,  
use TTL or policy-based methods.



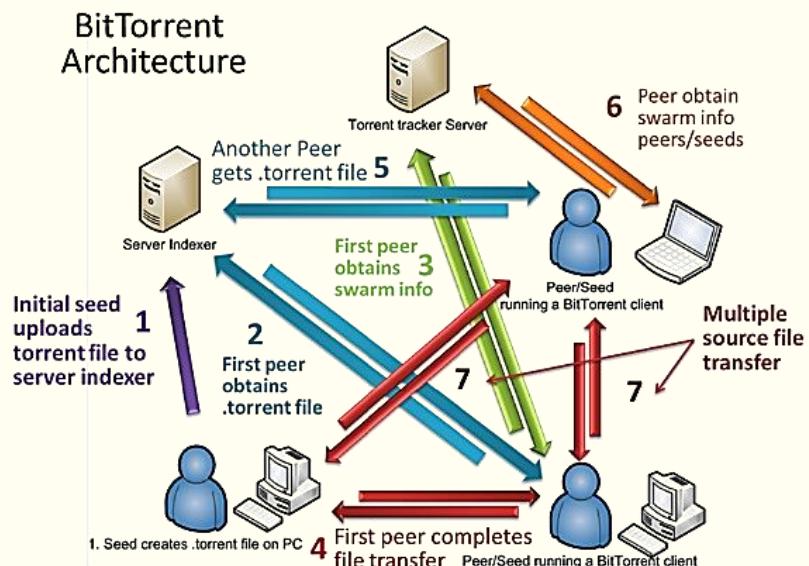
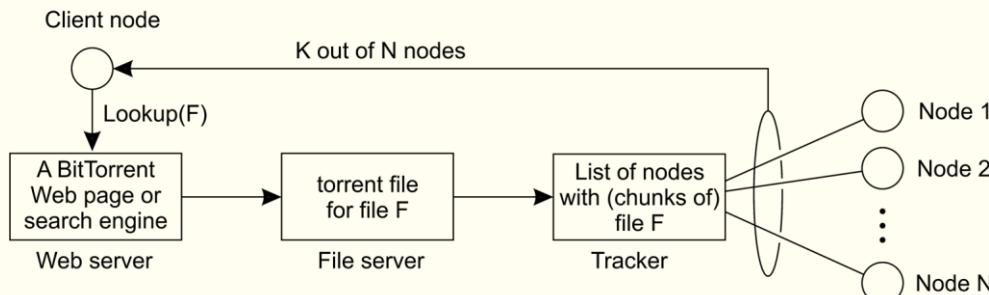
- Hierarchical P2P networks:

- A.k.a Super-peer overlay topologies.
- Sometimes, it is sensible to break the **symmetry** in P2P systems:
  - Having **index servers** improves the **performance** of searching in **unstructured P2P** systems, e.g., Content Delivery Network (**CDN**),
  - **Brokers** provide **efficient** data access.
- Super-nodes are always **fix** or **eligible** to be chosen by some **leader-election** technique.
- Weak peers can be **associated** with a **fixed** super-peer or **deterministically** or **randomly**.



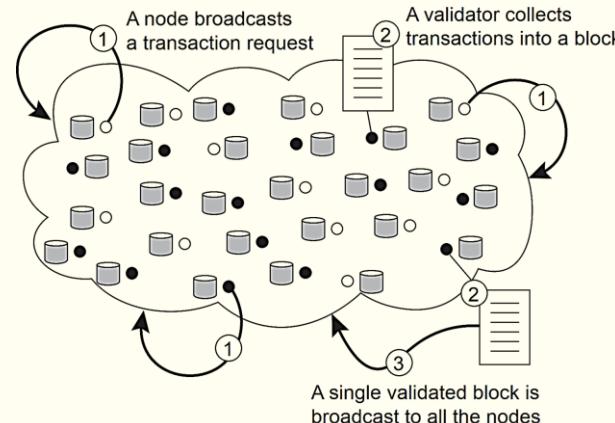
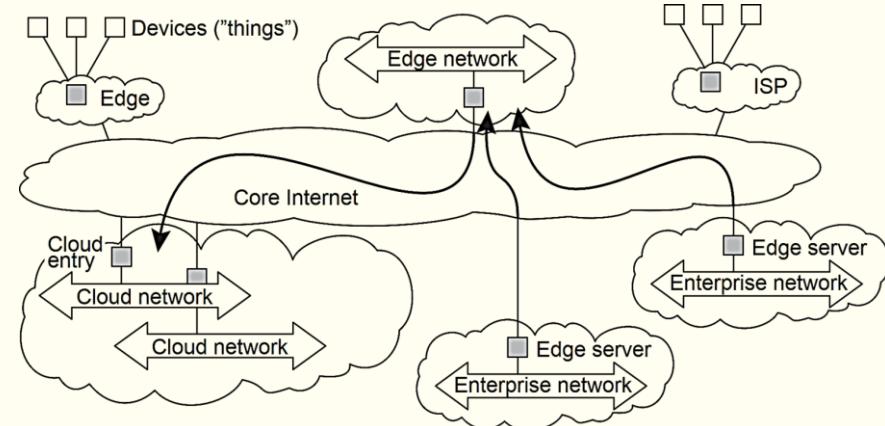
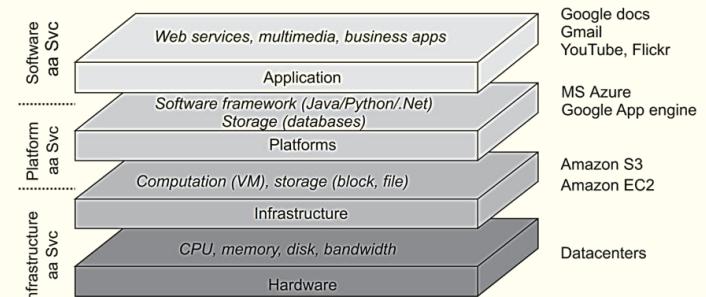
- Examples:

- BitTorrent.
- Skype



# Hybrid System Architectures:

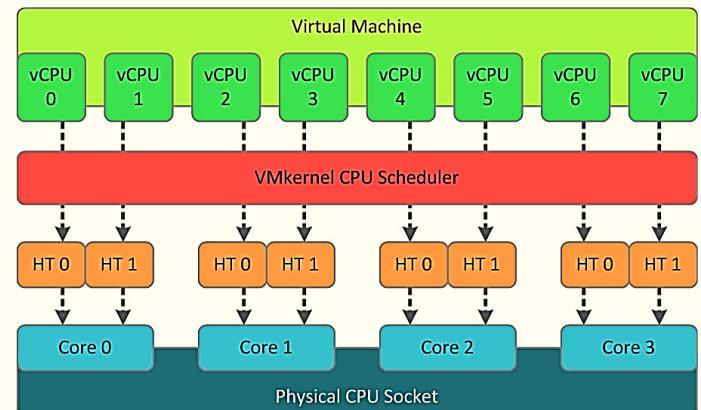
- Real-world **DSs** are complex, they combine hundreds of architectures to achieve the best features from all involved styles.
- Cloud Architecture:**
  - Provides configurable microservices across dedicated infrastructure,
  - Highly advanced **client/server** architecture with completely **hidden** servers (Function-as-a-Service (**FaaS**)).
- Edge-cloud Architecture:**
  - Servers are placed at the edge of the network (**Hierarchical Style**).
  - Better utilisation of latency & bandwidth.
  - Issues:** reliability, security, privacy, resource allocation, service placement, edge selection.
- Blockchain Architecture:**
  - Decentralised **P2P** distributed **ledgers**.
  - Validator requires distributed **consensus**:
    - Centralised!, distributed! (**permissioned**),
    - Decentralised! (**permissionless**).



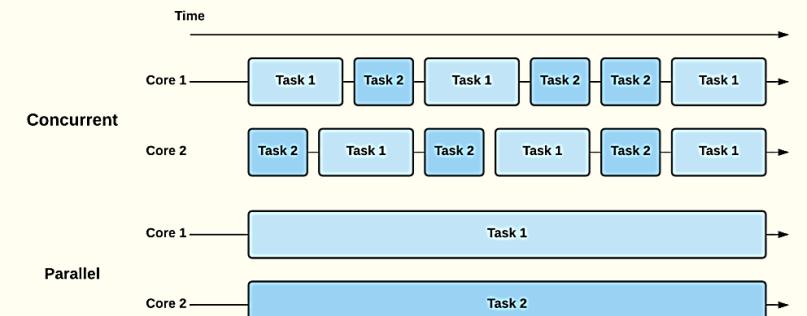
# PROCESSES

# Basics:

- Software **processes** originated in the **OSs**, they play crucial roles in **DSs**.
- In recent **OSs**, there are **virtual processors** built on top of **physical** processors:
  - **Multiprocessing**: processes are executed simultaneously in a multi-processors system.
  - **Multicore** is a processor that has more than one independent processing units (core).
  - **Multithreading**: threads of a process are executed simultaneously.
  - **Parallelism** is the ability to execute independent **processes** in the same instant of **time**.
  - **Concurrency** is to execute multiple **processes/threads** in an overlapping time period, i.e., not running at the same time.

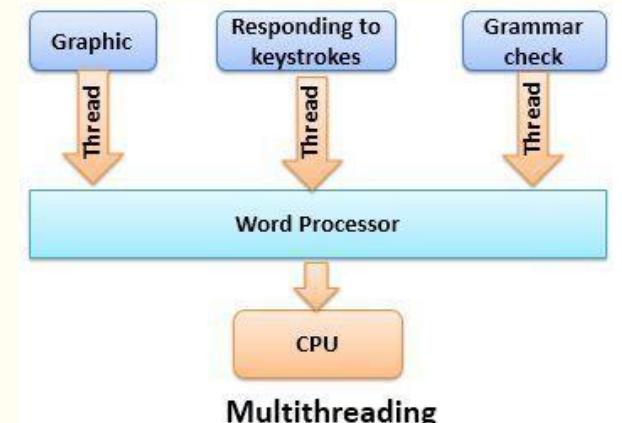
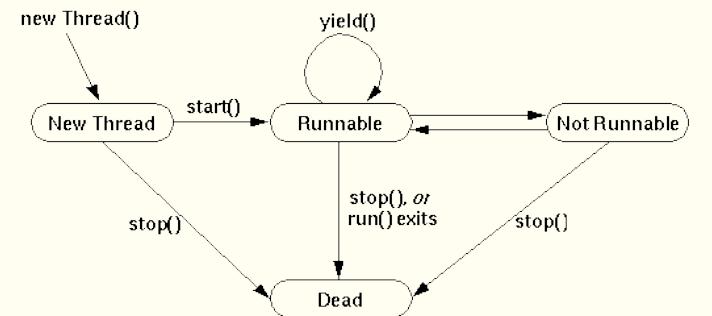
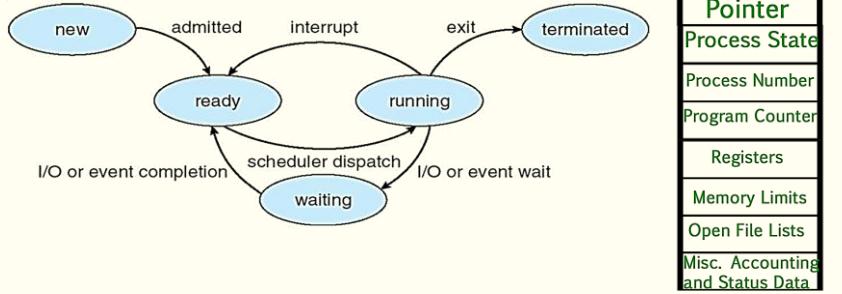


- This topic discusses many **parallel** and **concurrent computing** foundations:
  - Processes & Threads.
  - Zoom in Client/server Architecture.
  - Virtualisation.
  - Code Migration and Agents.



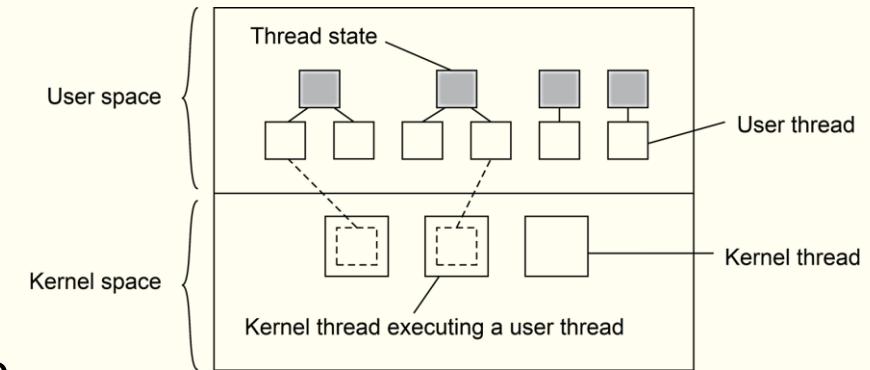
# Processes & Threads:

- An OS/DS process is a software **program** dispatched from the **ready state** and **scheduled** in the **CPU** for **execution** in an OS's virtual processor:
  - Each process has **independent address** space provided by the **OS**,
  - OSs identify** and **tracks** processes via the Process Control Block (**PCB**),
  - It is **heavyweight**, takes **time** to **create**, **terminate**, and **switch contexts!**
  - It is **isolated (HW protected)** and preserves **less efficient communication**.
- A thread is the **segment** (finer granularity) of a **process**, i.e., a process can contain **multiple threads**:
  - Threads are **cheaper to create, terminate and switch!**,
  - They **share the same address space**, not **isolated** (i.e., not protected against each other) and provide more efficient **communication**,
  - A **multithreaded application** performs highly **better** than its single-threaded counterpart.
- Why use threads?**
  - Avoid needless **blocking** in a multi-threaded process, e.g., when doing **I/O**.
  - Support **parallelism** as threads can **run in parallel** on a multi **core/processor**.
  - Avoid **process switching** as thread switching is faster, why?
- Discussion:** is an android/java thread similar in structure to a windows thread?

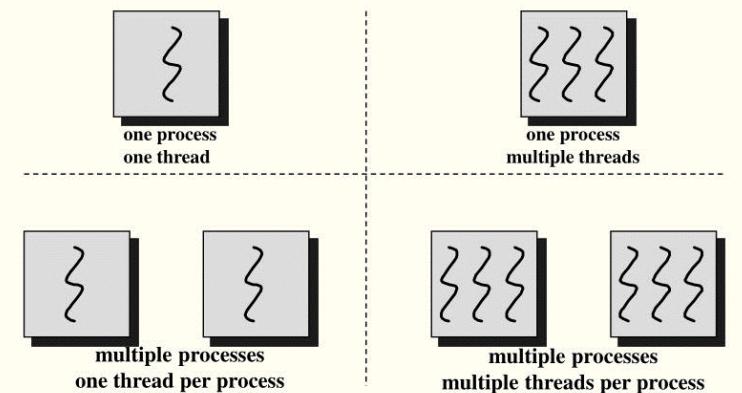


- Threads Implementation:

- Typically, threads are provided in programming package/library, which contains operations to create, destroy and synchronise threads (e.g., mutexes and condition variables).
  - User-level threads (many-to-one-model): implemented by the user, cheap switching but single schedulable entity.
    - A blocking system call blocks the entire process!
  - Kernel-level threads (one-to-one-model): implemented by OSs, transparent to the application.



- User-level and Kernel-level threads are combined into a single concept (Hybrid Threads), leading to performance gain but with increased complexity (Lightweight Processes).
- Kernel-level threads are runnable threads that execute User-level threads.
- In practice, applications are a collection of concurrent processes, and make use of the inter-process facilities in the OS, e.g., Web-Servers.
- Processes have isolated memory space and HW protected; but heavyweight.
- Threads are faster and lighter but difficult to implement as developers are responsible for managing concurrent access to data.



- Threads in DSs:

- Client-side uses multi-threaded computing with high degree of transparency, e.g., to:
  - Hide network latency in Web-browser (multiple downloads).
  - Make simultaneous calls each by a different thread, then each thread waits for the response.

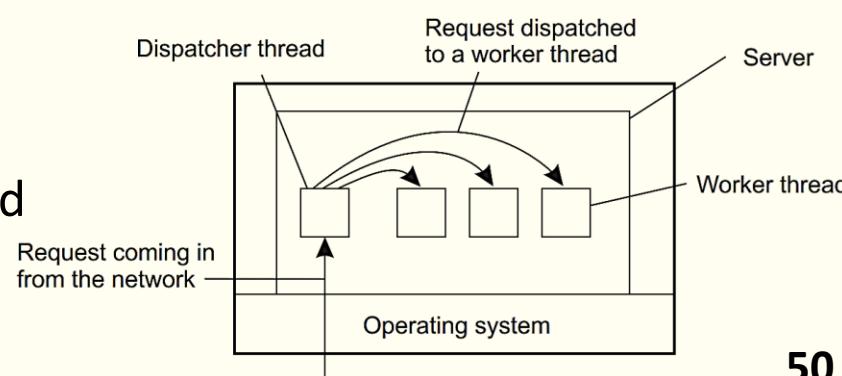
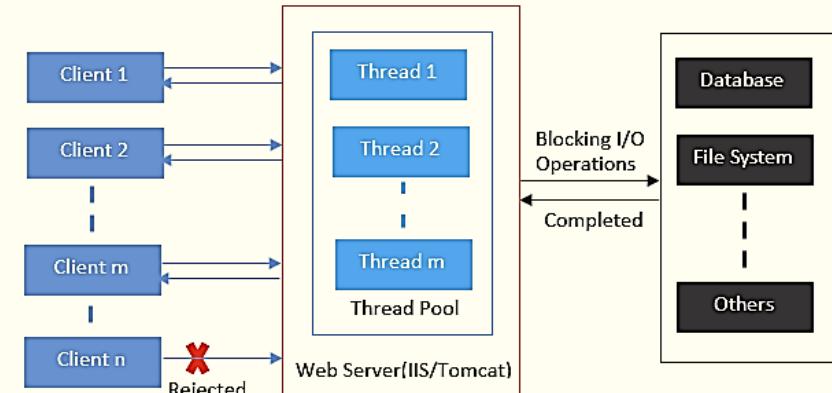
- Server-side use multi-threaded computing to achieve high performance, e.g.,

- By using cheap threads rather than processes.
- By hiding network latency with multithreaded connections.
- By using optimised I/O blocking calls.
- By adopting Dispatcher/Worker model
- Server construction options:

1. Single-threaded process: provide no parallelism and preserve blocking system calls,

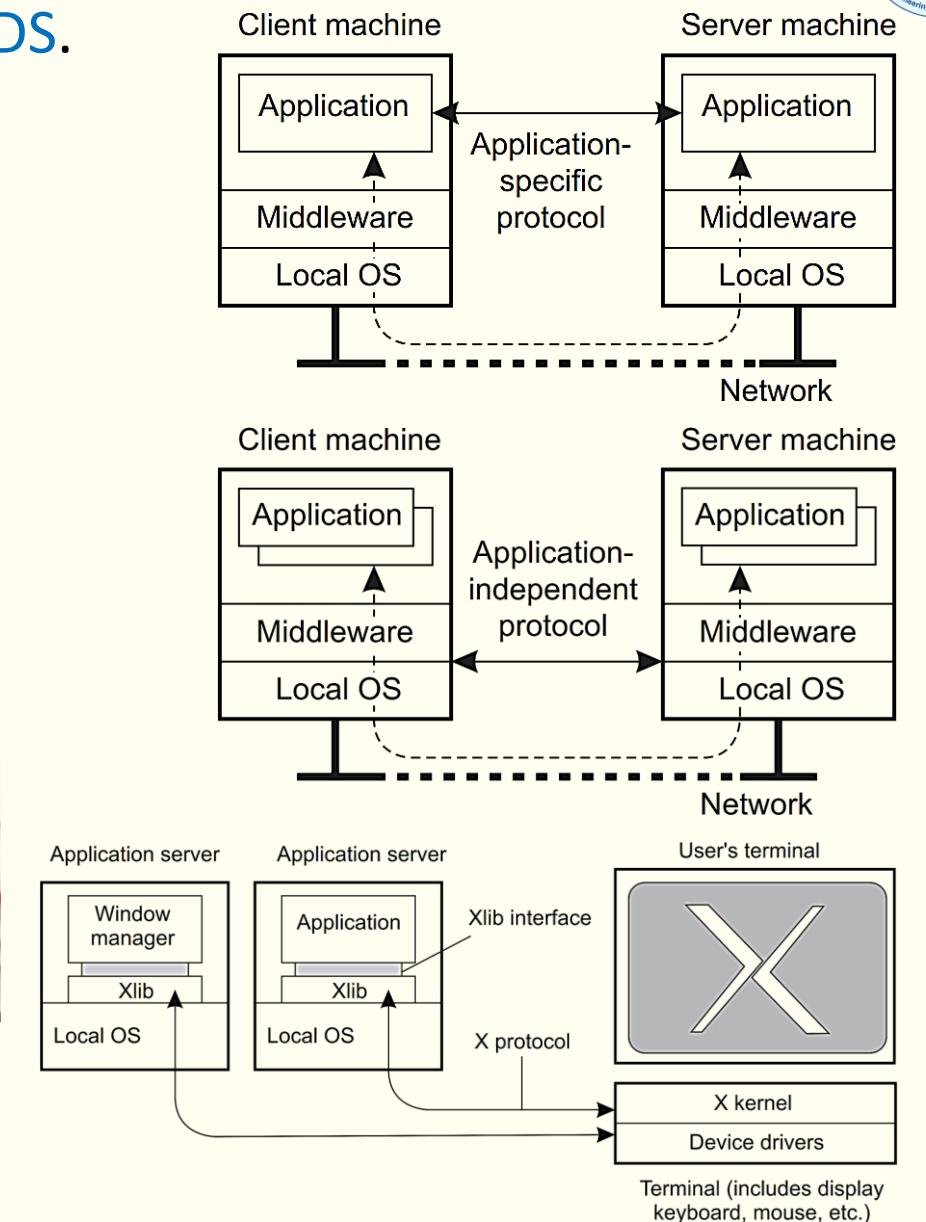
2. Threads (no supervisor): support parallelism with blocking system calls, i.e., in a single-processor machine, a thread may cause I/O or HW blocking.

3. Finite-state (with supervisor) machine: on an event/request, a dispatcher thread (supervisor) activates a thread to react; the dispatcher saves threads' state on the context-switching and switches context for block-causing threads; provides no parallelism and nonblocking system calls.



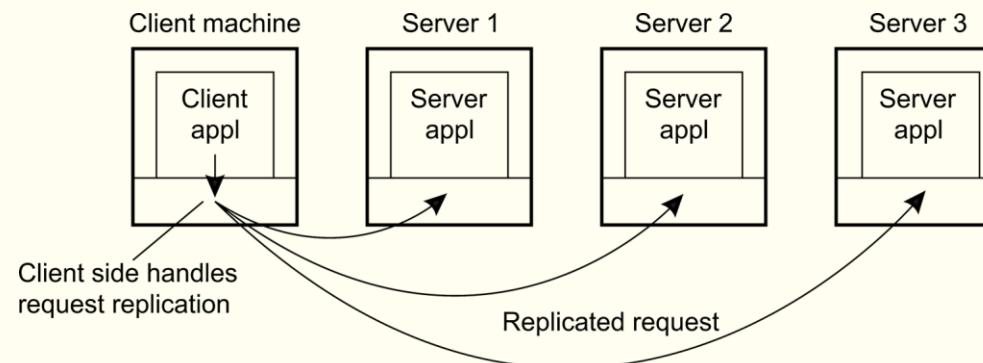
# Anatomy of Clients:

- Clients typically are means for users to interact with the DS.
- There are two interaction ways:
  1. A client machine has a separate counterpart/component that communicates with a remote service over a network, e.g., Calendar Apps.
    - Requires an application-level protocol to handle the communications/synchronisation,
  2. A client machine provides direct access to remote services by offering only a convenient user interface, e.g., command line terminals.
    - There is no local data storage or processing.
- Examples:
  - X-window system acts as a client to the X-kernel running as a server.
  - Virtual Network Computing (VNC).
  - Virtual Desktop Environment (e.g., RDP, AnyDesk, TeamViewer).
  - Web Browsers and Document Object Model (DOM).
  - Progressive Web Apps (PWA).



- Client-side software:

- Client software comprises more than user interfaces; some computing in a client-server application takes place on the client side.
- Generally tailored for distribution **transparency**:
  - **Access**, hiding data/resources representation and access; e.g., clients **RPC** stubs handle remote objects access,
  - **Location/migration /relocation**, hiding server/resources location; clients can track actual/new server/resources location via naming services and brokers,
  - **Replication**, hiding server/resources replicas; the client stub may issue multiple invocations to replicas,



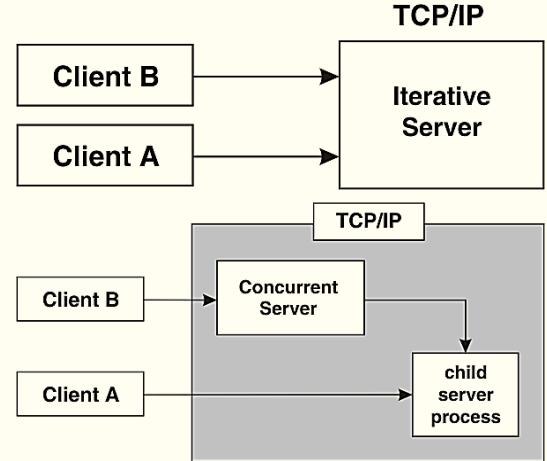
- **Failure**, through masking server/resources and communication failures; recovery via retransmission and multiple connection techniques.
- **Concurrency**, typically, transaction monitors distribute loads for load-balancing; enabling server/resource sharing among clients; require less support from client software.
- **Discussion**: what roles do **multithreading** & **virtualisation** play in client distribution **transparency**?

# Servers Organisation:

- Generally, servers **wait** (**Passive Process**) for an incoming **request** from a client, **process** the request, **wait** for the **next** incoming **request**.

- Basic server types:

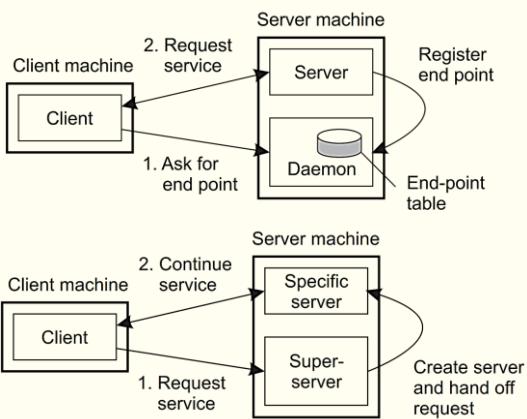
- Iterative**: a server handles the request before attending a next request.
- Concurrent**: via a **dispatcher** that picks up an incoming **request** that is then passed on to a separate **thread/process**.  
The **norm**, notably in the presence of blocking.



- Issues in contacting a server:

- Clients send requests to an **end-point (port)** where a service is listening.
  - How do clients **know** the **end-point** of a service?
    - One approach is to **globally assign end-points** for **well-known services**, typically, assigned by the Internet Assigned Numbers Authority (**IANA**).
    - Another approach is **dynamically assigned end-points** via a hidden service (**daemon**) that tracks services **end-points**. However, this approach is resource-draining.
    - More efficient approach is to implement a single **super-server** listening to all **end-points** associated with alive services

ftp-data	20	File Transfer [Default Data]
ftp	21	File Transfer [Control]
telnet	23	Telnet
smtp	25	Simple Mail Transfer
www	80	Web (HTTP)

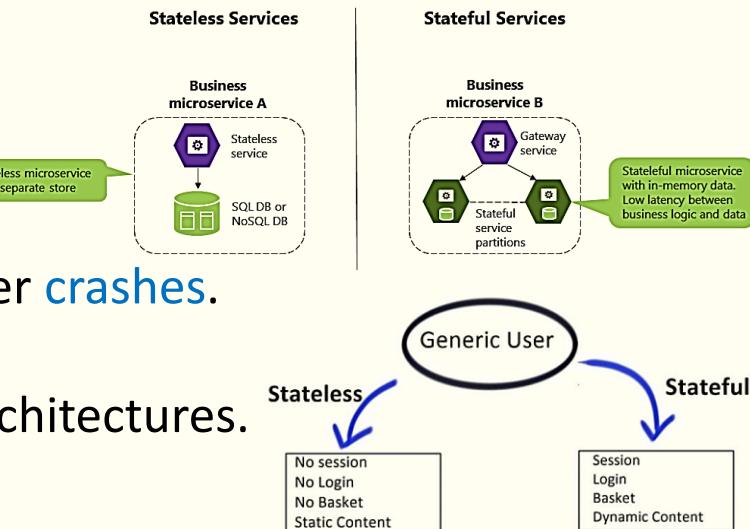


## 2. Interrupting a server:

- Is it possible to **interrupt** a server once it has **accepted** a service request?
- Examples: cancelling a file upload, or closing a playing stream.
- The server will **eventually** tear down the old connection, thinking the **client** has **crashed**.
- A better approach, use a separate **thread/process (port)** for **urgent** requests (**out-of-band** data),
- Another approach, use **facilities** of the **transport** layer, e.g., **TCP** and **OS** signalling techniques.

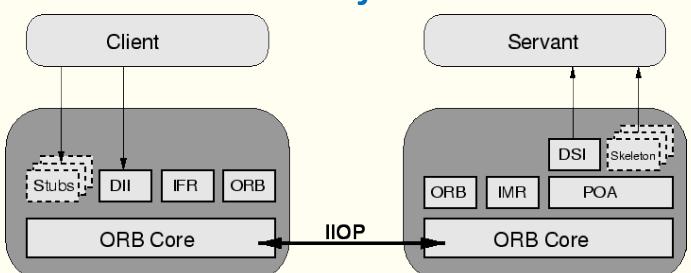
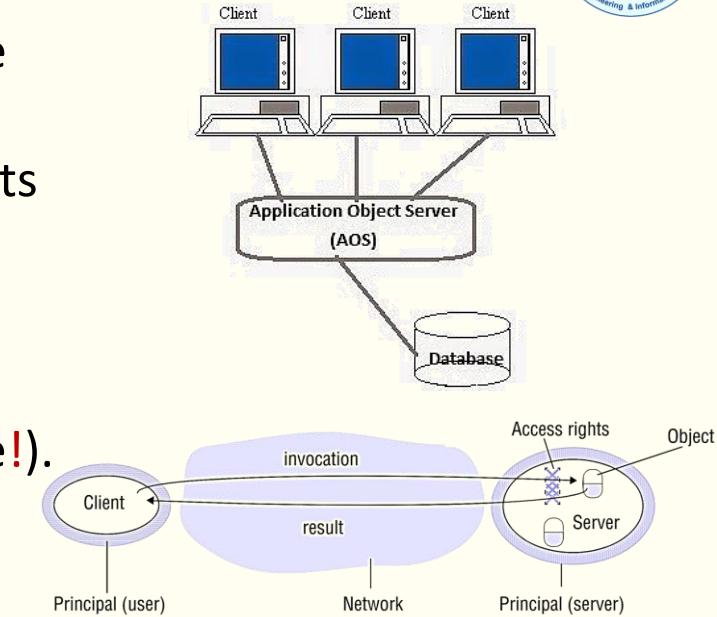
## 3. Stateless and stateful servers:

- A **stateless** server does not keep information about its clients, and can **change** its own **state** without informing clients, e.g., **HTTP** requests to Web servers. So: clients and servers are completely **independent** and have minimum state **inconsistencies after crashes**.
- A **soft-state** server **maintains state** on behalf of the client for a limited **time** (Session State),, e.g., maintained in three-tiered client-server architectures.
- A **stateful** server generally maintains **persistent information** (Operations, Crashes) on its clients. It improves **performance** as clients keep local **copies**, but preserve **complex recovery** on **failures** due to **recovering** all client's tracks.
- **States** in a stateful server can be **Sessional** or **Permanent**: vital information are stored in **DB**.
- The choice for a **stateless** or **stateful** design should not affect the **services** provided by the server, e.g., cookies and Web servers.
- **Discussion:** does connection-oriented communication fit into a stateless design?



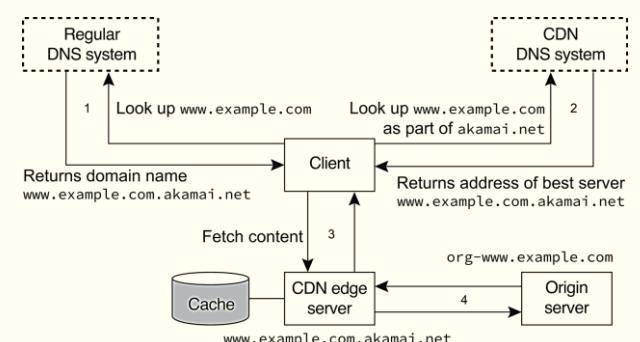
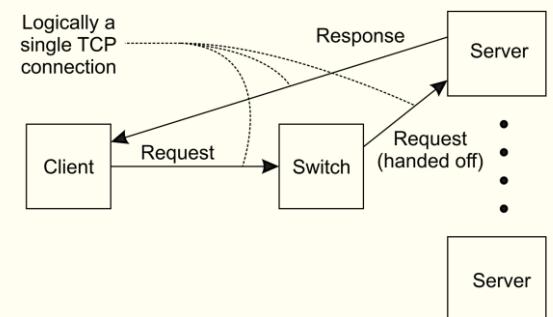
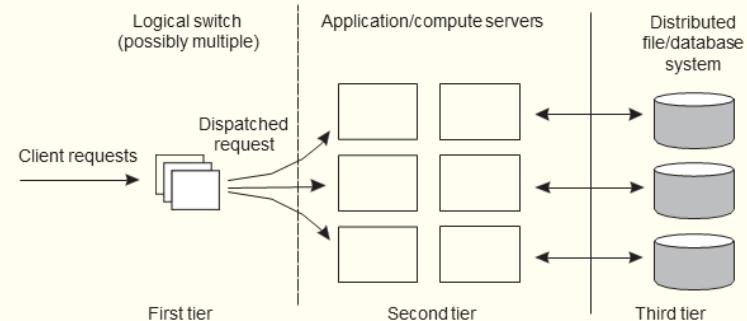
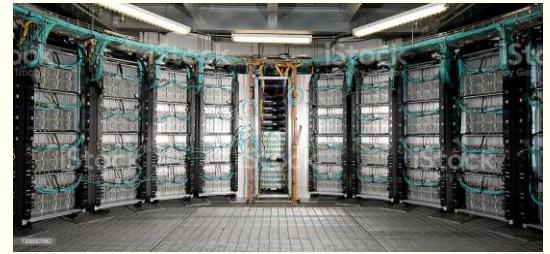
- Object Servers:

- Common **servers** are **object servers**.
- Object server by itself does not provide a specific service. Specific **services** are **implemented by the objects** that **reside in the server**.
- Alive remote object reside in a server consist of two parts: data representing its state and the code for executing its methods.
- Object server can either:
  - Keep callable objects **alive all the time** (resource-draining), or
  - **Activate callable objects** as long as they really **needed** (Performance Issue!).
  - Assign a **thread** for each callable **object**, or
  - Assign a **separate thread** for each **invocation request**.
- Object server use **generic** object **adapter/wrapper** to attach objects to a server application following a specific **activation policy**.
- An object **adapter** supports different **policies** via runtime configuration and can operate in **single-threaded** or **multithreaded** mode.
- A **servant** is the general **term** for a piece of **code** that forms the implementation of an **object**.
  - For examples of servers supporting activation policies, see [CORBA-compliant](#) systems.
  - For examples of servers that balance the separation between policies and mechanisms, study the [Apache Web server](#).



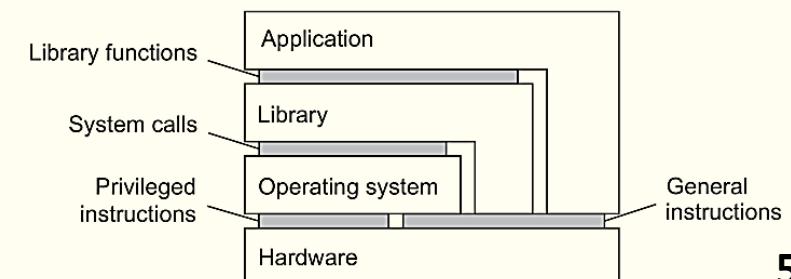
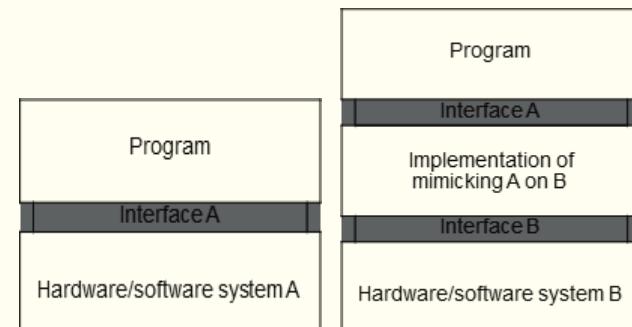
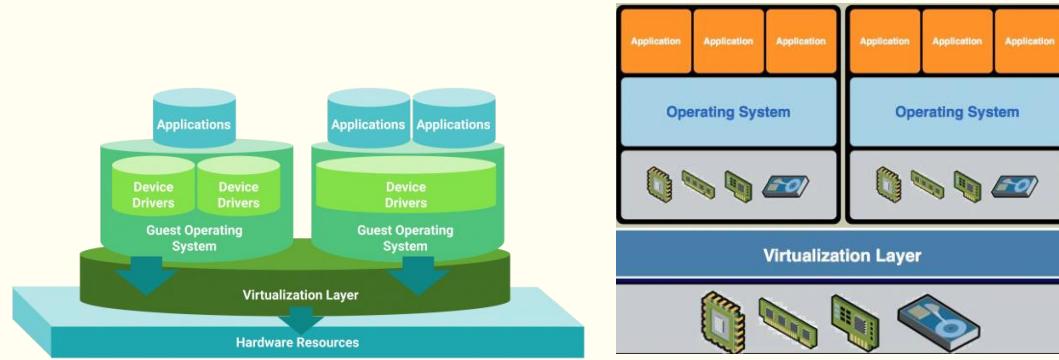
## • Server Clusters:

- A collection of **machines** connected through a **network**, where each machine **runs** one or more **servers**.
- Often, a server cluster is logically organized into **three-tiers**, not all clusters are **3-tiers**, e.g., media streaming are **2-tiers**.
- The **first tier** is responsible for passing **requests** to an appropriate server: **request dispatching** (Access Transparency).
- Access transparency is maintained by address/content request switching:
  - Transport-layer switches (**TCP**), e.g., **NAT**, **IPv6**, or
  - Application-layer switches, e.g., **URL**.
- Wide-area server clusters (Grid Servers) aim to offer data and services that are close to clients, e.g., Amazon, Google, **PalnetLab**, and Content Delivery Networks (**CDN**).
- For **scalability** and **availability**, a cluster may have **multiple access points**,
- **Origin server** maintains a **home address**, but **hands off** connections to **address** of collaborating **peer**, **origin server** and **peer** appear as one **server**.



# Virtualization:

- Virtualization is applied since 1990s, but has received attention due to the recent DSs.
  - To handle hardware changes that is faster than software,
  - To ease of portability and code migration,
  - To help in isolation of failing or attacked components.
  - It is vital concept for DSs, it helps achieve portability and handle failures
- In recent OSs, there are virtual processors built on top of physical processors:
  - A native service of an OS that allows an application/program to run concurrently with other applications, but highly independent of the underlying hardware and platforms.
- In practice, every DS offers a programming interface to higher-levels:
  - Instruction Set Architecture (ISA): between hardware and software, with two subsets:
    - Privileged instructions: allowed to be executed only by the OS.
    - General instructions: can be executed by any program.
  - System calls as offered by an operating system.
  - Library calls, a.k.a Application Programming Interface (API).



• Virtualization can take place in different ways:

- What examples for each scenario (a),(b), and (c)?
- Common issues: portability, transparency, performance, fault-tolerance, and security,
- Main issue is the implementation of control-sensitive and behaviour-sensitive instructions,  
Solution is modify guest OS (**Paravirtualization**).

• Containers:

- They are virtual machines run applications **side-by-side** but **independently** across different **platforms**.
- They require significant efforts to ensure **portability** and **performance**.
- A container is a collection of **binaries**/platform **images** that constitute the environment for running applications.
- Applications operating in different containers need to be isolated.
- Copying an entire platform is not efficient.
- Hosting platform need some control over containers.
- Very common in **Linux** platforms.
- Example: [PlanetLab](#) is a **DS** system for collaborative research.

• Discussion: what is better: virtual machines or containers?

The diagram illustrates three models of virtualization:

- (A) Full virtualization:** Application/Libraries run on top of a Runtime system, which runs on an Operating system, which runs on Hardware. The hardware interface is simulated.
- (B) Paravirtualization:** Application/Libraries run on top of an Operating system, which runs on a Virtual machine monitor, which runs on Hardware. The hardware interface is modified.
- (C) Containerization:** Application/Libraries run directly on an Operating system, which runs on Infrastructure (Host Operating System). This model is shown for both Hypervisor-based (Guest OS) and Docker Engine-based (Docker Engine) approaches.

The diagram shows a layered architecture of a Linux enhanced operating system:

- User-assigned virtual machines:** Multiple Vserver instances (each with /dev, /usr, /home, /proc) run on separate Linux enhanced operating systems, which are managed by a single underlying Linux enhanced operating system.
- Privileged management virtual machines:** A single Vserver instance (with /dev, /usr, /home, /proc) manages multiple User-assigned virtual machines.
- Hardware:** The bottom layer.

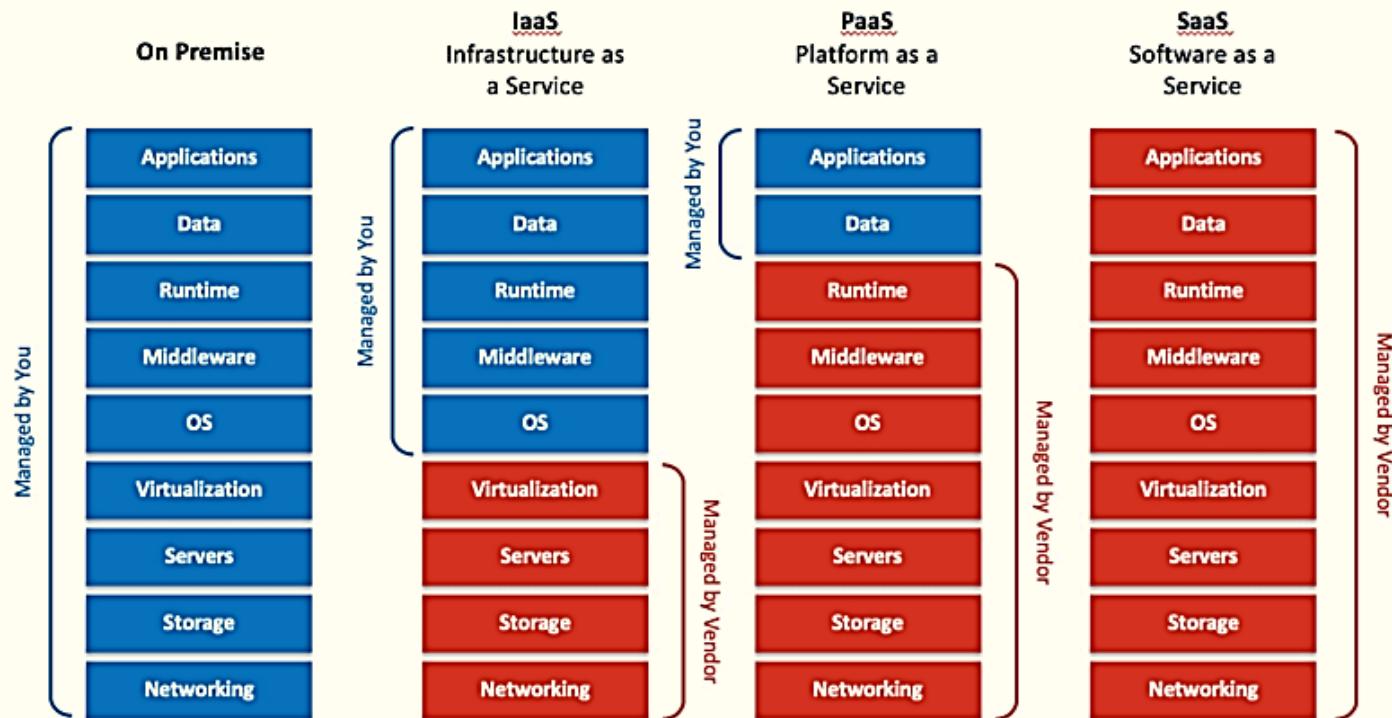
Al-Azhar University-Gaza  
Faculty of Engineering & Information Technology

58

M. M. Ayiad, Copyrights © 2023 All Rights Reserved, Al Azhar University-Gaza, Palestine.

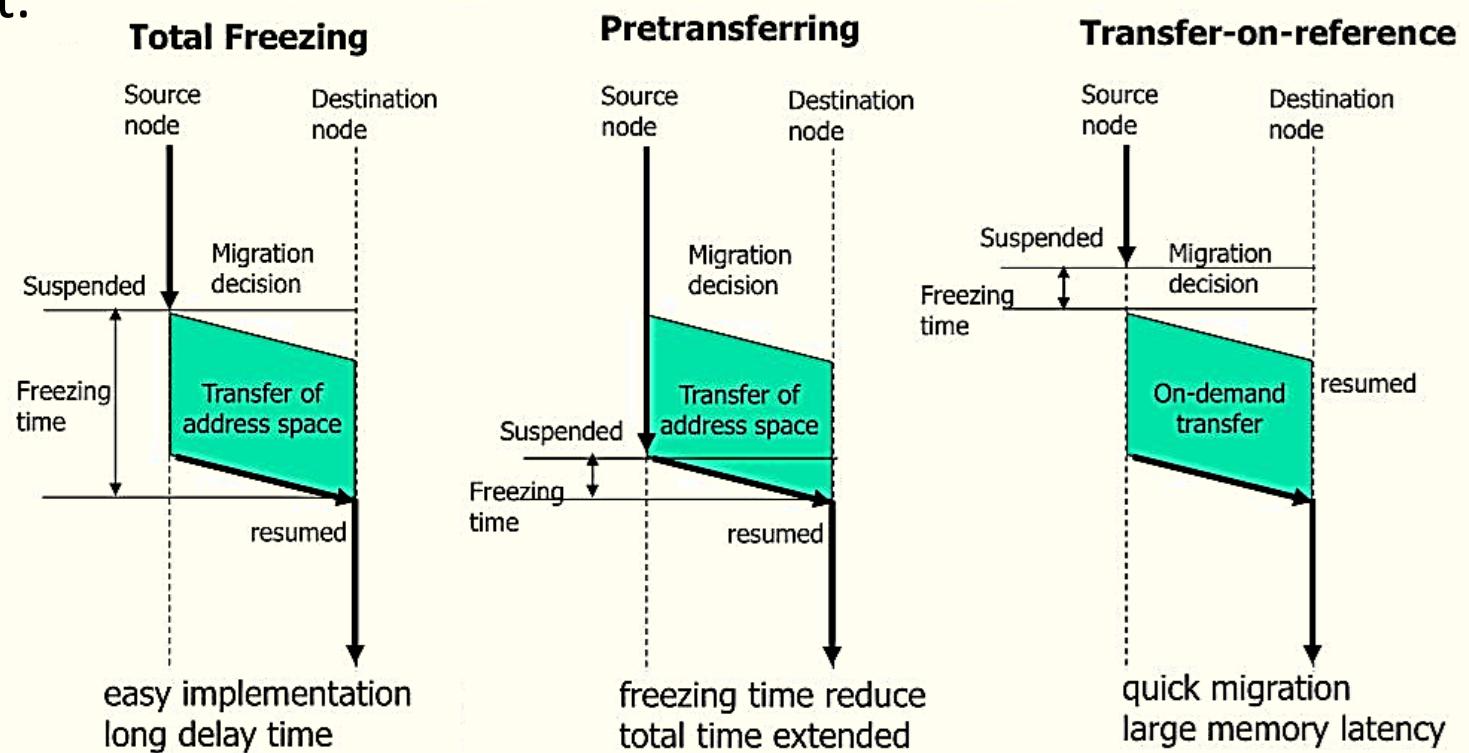
- Application of virtual machines to DS (Cloud Computing):

- A cloud provider rents-out a **virtual machine** and shares a physical machine among customers.
- A cloud provider ensures virtual machines' **isolation** and **performance**.
- Clouds provide three types of services:
  - Infrastructure-as-a-Service (IaaS)** covering the basic infrastructure,
  - Platform-as-a-Service (PaaS)** covering system-level services,
  - Software-as-a-Service (SaaS)** containing actual applications.



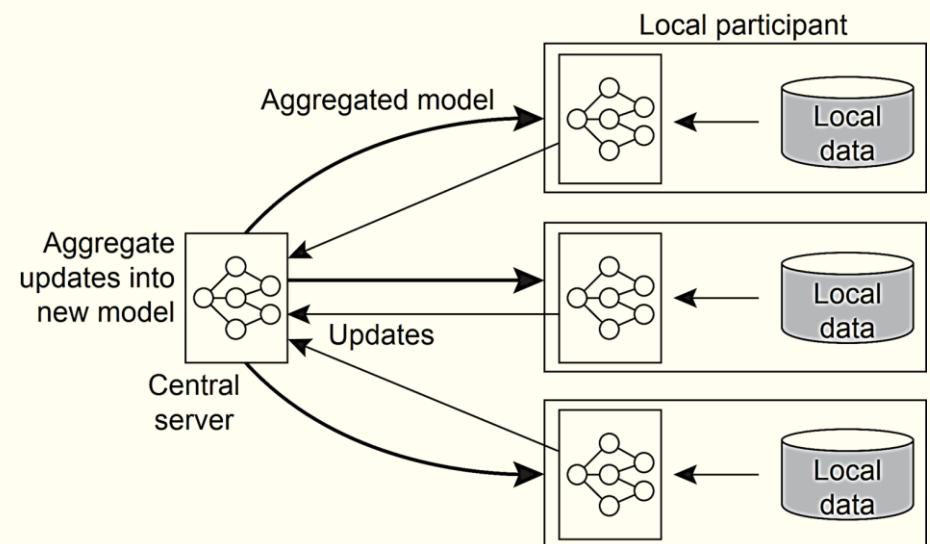
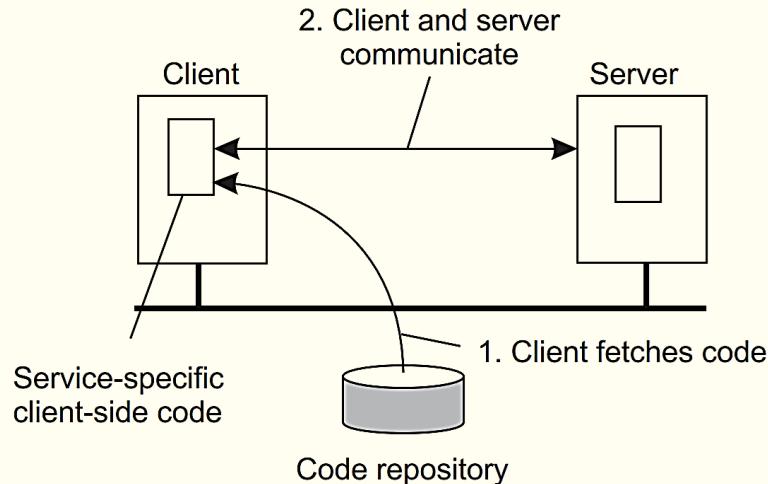
# Code Migration:

- So far, distributed **objects/components/services** communicate to **pass data**.
- There are situations in which **passing programs/codes** simplifies the **DSs**.
- Historically, code **migration** meant **process migration**, i.e., An entire process was moved from a node to another.
- In **DSs**, optimisation of computing power is a less pressing issue than minimising communication cost.



- Why migrate code?

- Often **migration decision** is based on **qualitative** reasoning instead of **mathematical** models, e.g. migrate **DB queries** to the server rather than sending them over the network.
- **Performance**: move processes to **lightly loaded!** nodes (i.e., **offloading**), Also, via parallelism and agents, e.g., search mobile agents.
- **Load distribution**:
  - To ensure that servers in a data-centre are sufficiently loaded, e.g., to save energy,
  - To minimize communication by that computations are close to where the data is, e.g., data training & mining (**federate learning**).
- **Privacy and security**: move process where data is for security.  
**Flexibility**: migrate code to a **client** when **needed**, e.g., moving an interface to access server's file system.

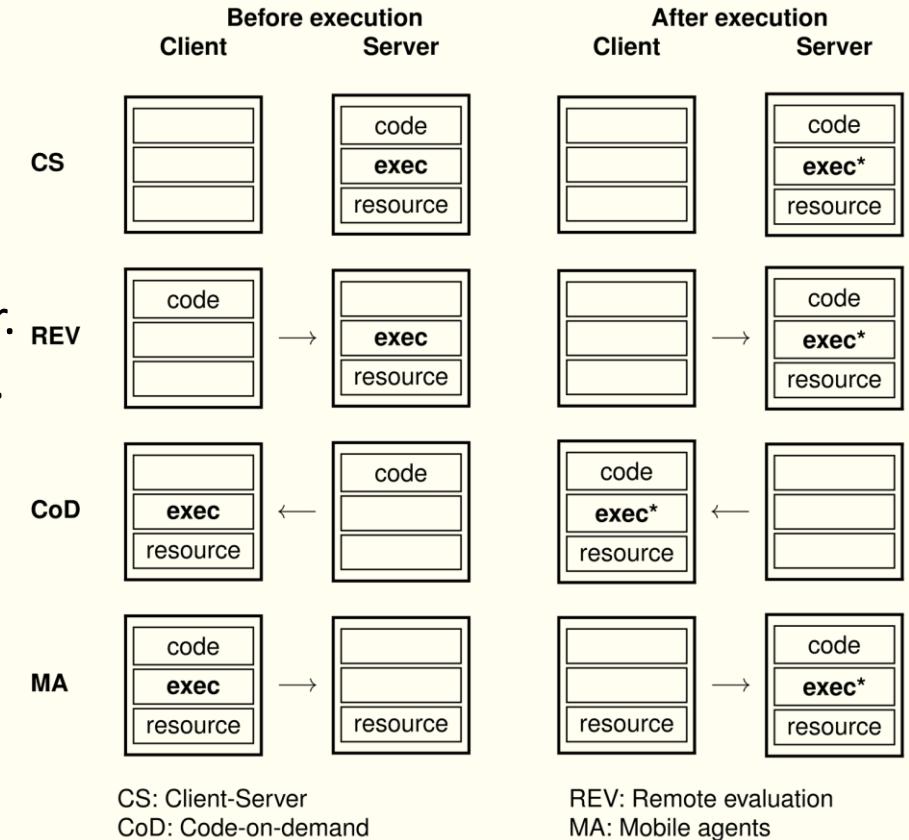


- **Models for code migration:**

- To understand the model for code migration, refer to a framework proposed by [Fuggetta et al.](#)
- **Code segment** contains the set of instructions of the program being executed.
- **Resource segment** contains references to external resources needed by the process, e.g., files, printers.
- **Execution segment** is used to store the current execution state of a process.
- **Sender-initiated** migration.
- **Receiver-initiated** migration.

- The four different **paradigms** for code **mobility**,

- Client-server (**CS**): segments are in and run at the server.
- Remote Evaluation (**REV**): code migrates for execution at the sever.
- Code-on-Demand (**CoD**): code migrates for execution at the client.
- Mobile Agents (**MA**): segments migrate to server for execution.
  - **Weak** mobility: migrate **code**, **data**, and **reboot** execution.
  - **Strong** mobility: migrate **code**, **data**, and **continue** execution.
    - Either by **migrating** the whole object or by **cloning** it.



- For code migration, the problems coming from **heterogeneity!**  
They are in many respects the same as those of **portability**.

# Review Questions:

1. Resources sharing is an objective of **DSS**, give examples of hardware, data, and software resources that might be shared in a **DS**. Explain the most needed requirements for efficient resource sharing.
2. Discuss two of the following **DSs** challenges: heterogeneity, openness, fault handling, concurrency, quality of service, scalability, and transparency.
3. Clock synchronisation is problematic in **DSS**s, explain the factors that limit the accuracy of the clocks. How could organise a global clock in a large system? Discuss the accuracy of global clock synchronisation mechanisms .
4. What is novel about cloud computing as a concept in comparison to traditional client-server computing?
5. A server program written in C++ provides objects that are intended to be accessed by clients. Clients are written in Java. Both server and client objects have access to the internet. Describe the problems due to heterogeneity that need to be solved to make it possible for a client object to invoke a method on a server object.
6. Discuss the sharing, scalability, and anonymity in different peer-to-peer architectures.
7. Discuss the query and data dissemination in the peer-to-peer architectures. Consider centralised, structured, unstructured, and hierachal topologies.
8. Computing components may fail when a client process invokes a method in a server object, discuss with examples the expected failures. Suggest how components can be made to tolerate such failures.
9. Resources in the **WWW** and other services are named by **URLs**. Discuss? Give examples of web resources that can be named by **URLs**.
10. What problems are foreseen in the direct coupling between communicating entities that is implicit in remote invocation approaches? Consequently, what advantages can be anticipated from a level of decoupling as offered by referential and time uncoupling?
11. List one typical criterion for selecting a node to be promoted to a super node. Justify why such a choice would improve network stability.
12. Choose a major Internet application of your favourite, and discuss with illustrations the client-server architecture of that application..
13. A search engine is a web server that responds to client requests to search in its stored indexes and concurrently runs several web crawler tasks to build and update the indexes. discuss and explain the requirements for synchronisation between these concurrent activities.
14. Most peer-to-peer systems are often simply personal computers in users' offices or homes. What are the implications of this for the availability and security of any shared data objects that they hold and to what extent can any weaknesses be overcome through the use of replication? .
15. Assume a transaction processing system that maintains ACID. Briefly describe one serious shortcoming of each of the following implementations: (a) The database is updated on disk with each transaction. (b) The database is kept in memory and on disk, with the copy on disk updated every 50 transactions. (c) The database is kept in memory. A log file is maintained on disk recording every transaction.
16. Discuss the Middleware architecture styles.
17. Consider any business company. Sketch out a three-tier solution to the distributed service of the company. Use this to illustrate the benefits and drawbacks of a three-tier solution considering issues such as performance, scalability, dealing with failure and also maintaining the software over time.
18. For each of the factors that contribute to the time taken to transmit a message between two processes over a communication channel, state what measures would be needed to set a bound on its contribution to the total time. Why are these measures not provided in current general-purpose distributed systems?
19. What is a case where it is preferable to use multiple threads instead of multiple processes? What is a case where the reverse is true?
20. Suppose you have a stateful service. What techniques could you use to scale it vertically or horizontally?

# COMMUNICATIONS